


# Neural Networks II: Backpropagation

---


Robin Jia  
USC CSCI 467, Spring 2025  
February 18, 2025

# Review: Neural Networks (2-layer MLP)



- Hidden layer = A bunch of logistic regression classifiers
  - Parameters:  $w_j$  and  $b_j$  for each classifier, for each  $j=1, \dots, h$
  - $h$  = number of neurons in hidden layer ("hidden nodes")
  - Produces "activations" = learned feature vector
- Final layer = linear model
  - For regression: linear model with weight vector  $v$  and bias  $c$

# Review: Neural Networks (2-layer MLP)




- Hidden layer = A bunch of logistic regression classifiers
  - Parameters:  $w_j$  and  $b_j$  for each classifier, for each  $j=1, \dots, h$
  - **Equivalently: matrix  $W$  ( $h \times d$ ) and vector  $b$  (length  $h$ )**
  - $h$  = number of neurons in hidden layer (“hidden nodes”)
  - Produces “activations”= learned feature vector
- Final layer = linear model
  - For regression: linear model with weight vector  $v$  and bias  $c$
- **Parameters of model are  $\theta = (W, b, v, c)$**

# Review: Neural Network Building Blocks

---

## (1) Linear Layer

- Input  $x$ : Vector of dimension  $d_{in}$
- Output  $y$ : Vector of dimension  $d_{out}$
- Formula:  $y = Wx + b$
- Parameters
  - $W$ :  $d_{out} \times d_{in}$  matrix
  - $b$ :  $d_{out}$  vector
- In pytorch: `nn.Linear()`




# Review: Neural Network Building Blocks

---

## (2) Non-linearity Layer


- Input  $x$ : Any number/vector/matrix
- Output  $y$ : Number/vector/matrix of same shape
- Possible formulas:
  - Sigmoid:  $y = \sigma(x)$ , elementwise
  - Tanh:  $y = \tanh(x)$ , elementwise
  - Relu:  $y = \max(x, 0)$ , elementwise
- Parameters: None
- In pytorch: `torch.sigmoid()`, `nn.functional.relu()`, etc.



# Review: Neural Network Building Blocks

## (3) Loss Layer

- Inputs:
  - $y_{\text{pred}}$ : shape depends on task
  - $y_{\text{true}}$ : scalar (e.g., correct regression value or class index)
- Output  $z$ : scalar
- Possible formulas:
  - Squared loss:  $y_{\text{pred}}$  is scalar,  $z = (y_{\text{pred}} - y_{\text{true}})^2$
  - Softmax regression loss:  $y_{\text{pred}}$  is vector of length  $C$ ,
$$z = - \left( y_{\text{pred}}[y_{\text{true}}] - \log \sum_{i=1}^C \exp(y_{\text{pred}}[i]) \right)$$
- Parameters: None
- In pytorch: `nn.MSELoss()`, `nn.CrossEntropyLoss()`, etc.



# Review: Training Neural Networks

---

## Linear Regression

- Model's output is

$$g(x) = w^\top x + b$$

- (Unregularized) loss function is

$$\frac{1}{n} \sum_{i=1}^n (g(x^{(i)}) - y^{(i)})^2$$

## Regression w/ Neural Networks

- Model's output is

$$g(x) = v^\top \sigma(Wx + b) + c$$

- **Use same loss function**, in terms of  $g$ !

$$\frac{1}{n} \sum_{i=1}^n (g(x^{(i)}) - y^{(i)})^2$$

## Training objective for both types of models:

$$\frac{1}{n} \sum_{i=1}^n \ell(y^{(i)}, g(x^{(i)})), \text{ where } \ell(y, u) = (y - u)^2$$

Also applies for  
logistic regression,  
softmax regression, etc.

# Review: Training Neural Networks

---

General loss function:  $\frac{1}{n} \sum_{i=1}^n \ell \left( y^{(i)}, g(x^{(i)}) \right)$

Model's output, depends on all model parameters  $\theta$  (includes all layers)

- How to minimize? Gradient Descent!

$$\theta \leftarrow \theta - \eta \cdot \underbrace{\frac{1}{n} \sum_{i=1}^n \nabla_{\theta} \ell \left( y^{(i)}, g(x^{(i)}) \right)}_{\text{Average of per-example gradients}}$$

Average of per-example gradients

- **Today: How to compute gradient of loss w.r.t. parameters for any neural network**
  - So many different ways to assemble building blocks
  - Don't want to re-do gradient calculations by hand each time
  - **Can we write an algorithm to do it?**



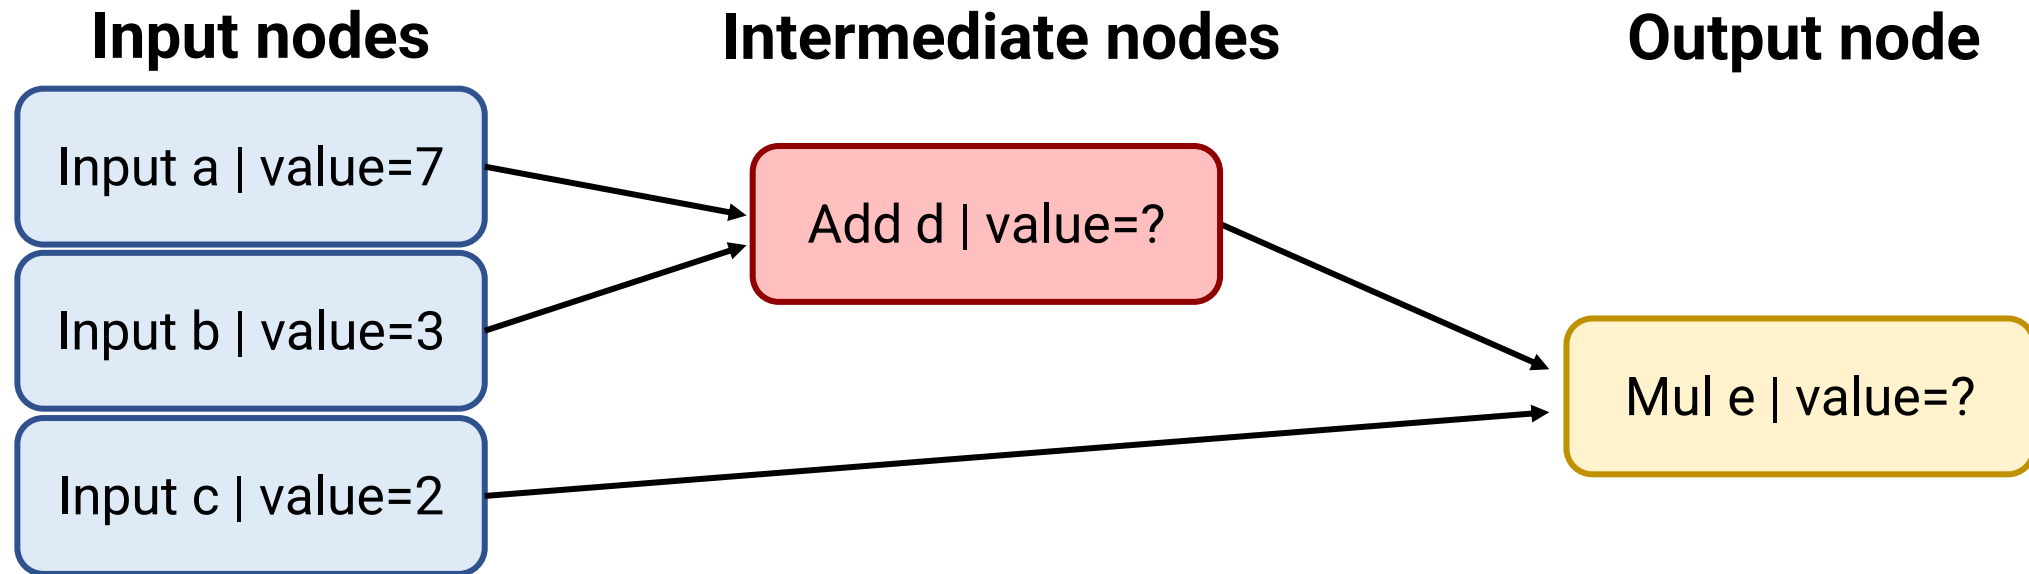
# Today's Plan

---

- The computation graph
- Backpropagation on trees
- Backpropagation on DAGs

# Computation Graph

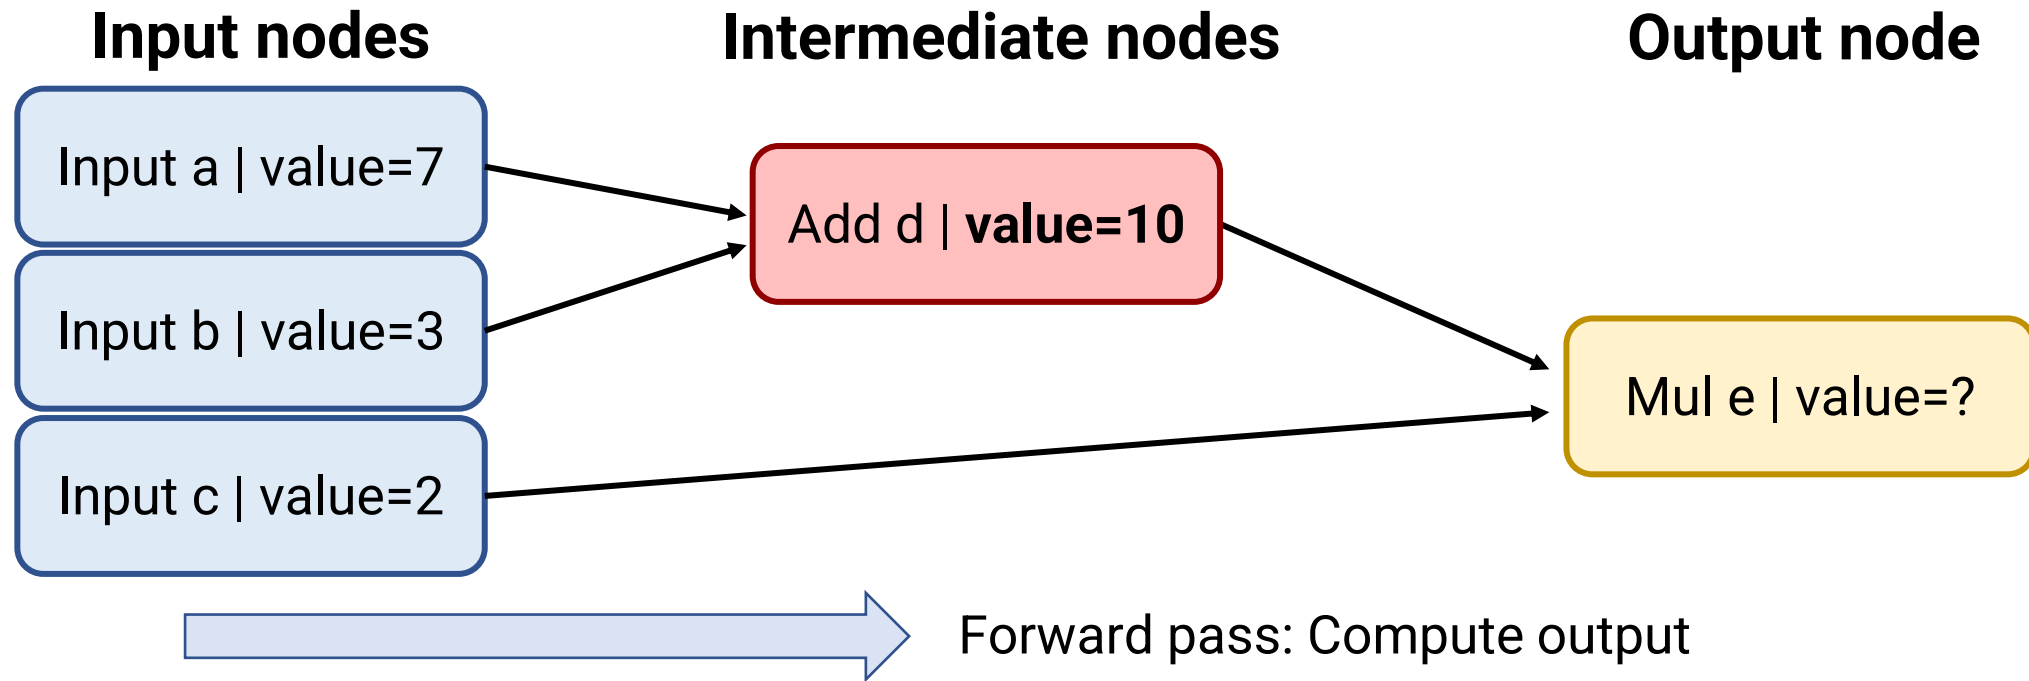
Computation graph for  $(a + b) * c$  when  $a=7, b=3, c=2$



Different way of drawing the “building blocks” of neural networks

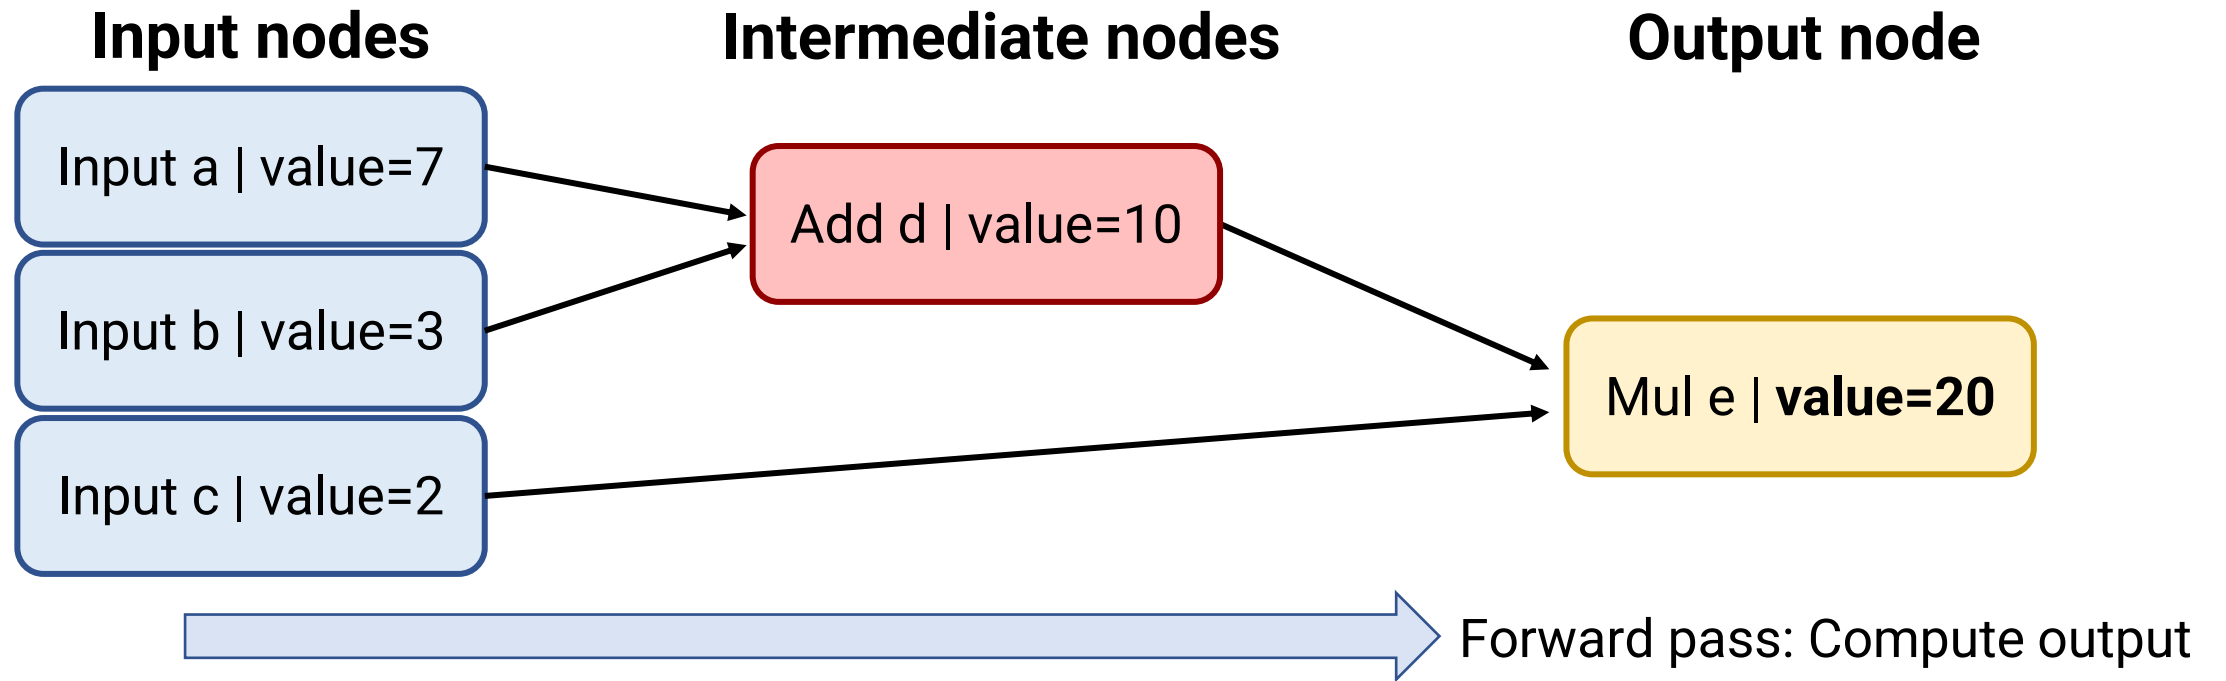
# Computation Graph

Computation graph for  $(a + b) * c$  when  $a=7, b=3, c=2$




# Computation Graph

Computation graph for  $(a + b) * c$  when  $a=7, b=3, c=2$



# Gradient checking

- **Numerical gradients:** A simpler but less efficient way to compute gradients
- What does  $\partial y / \partial x$  mean?
  - If I change  $x$  by epsilon, by what proportion of epsilon does  $y$  change?
- We can just compute this for every input node!
- Pro: Easy to implement, useful to check correctness
- Con: Slow—requires  $O(\text{\#inputs})$  function evaluations



# Let's implement!



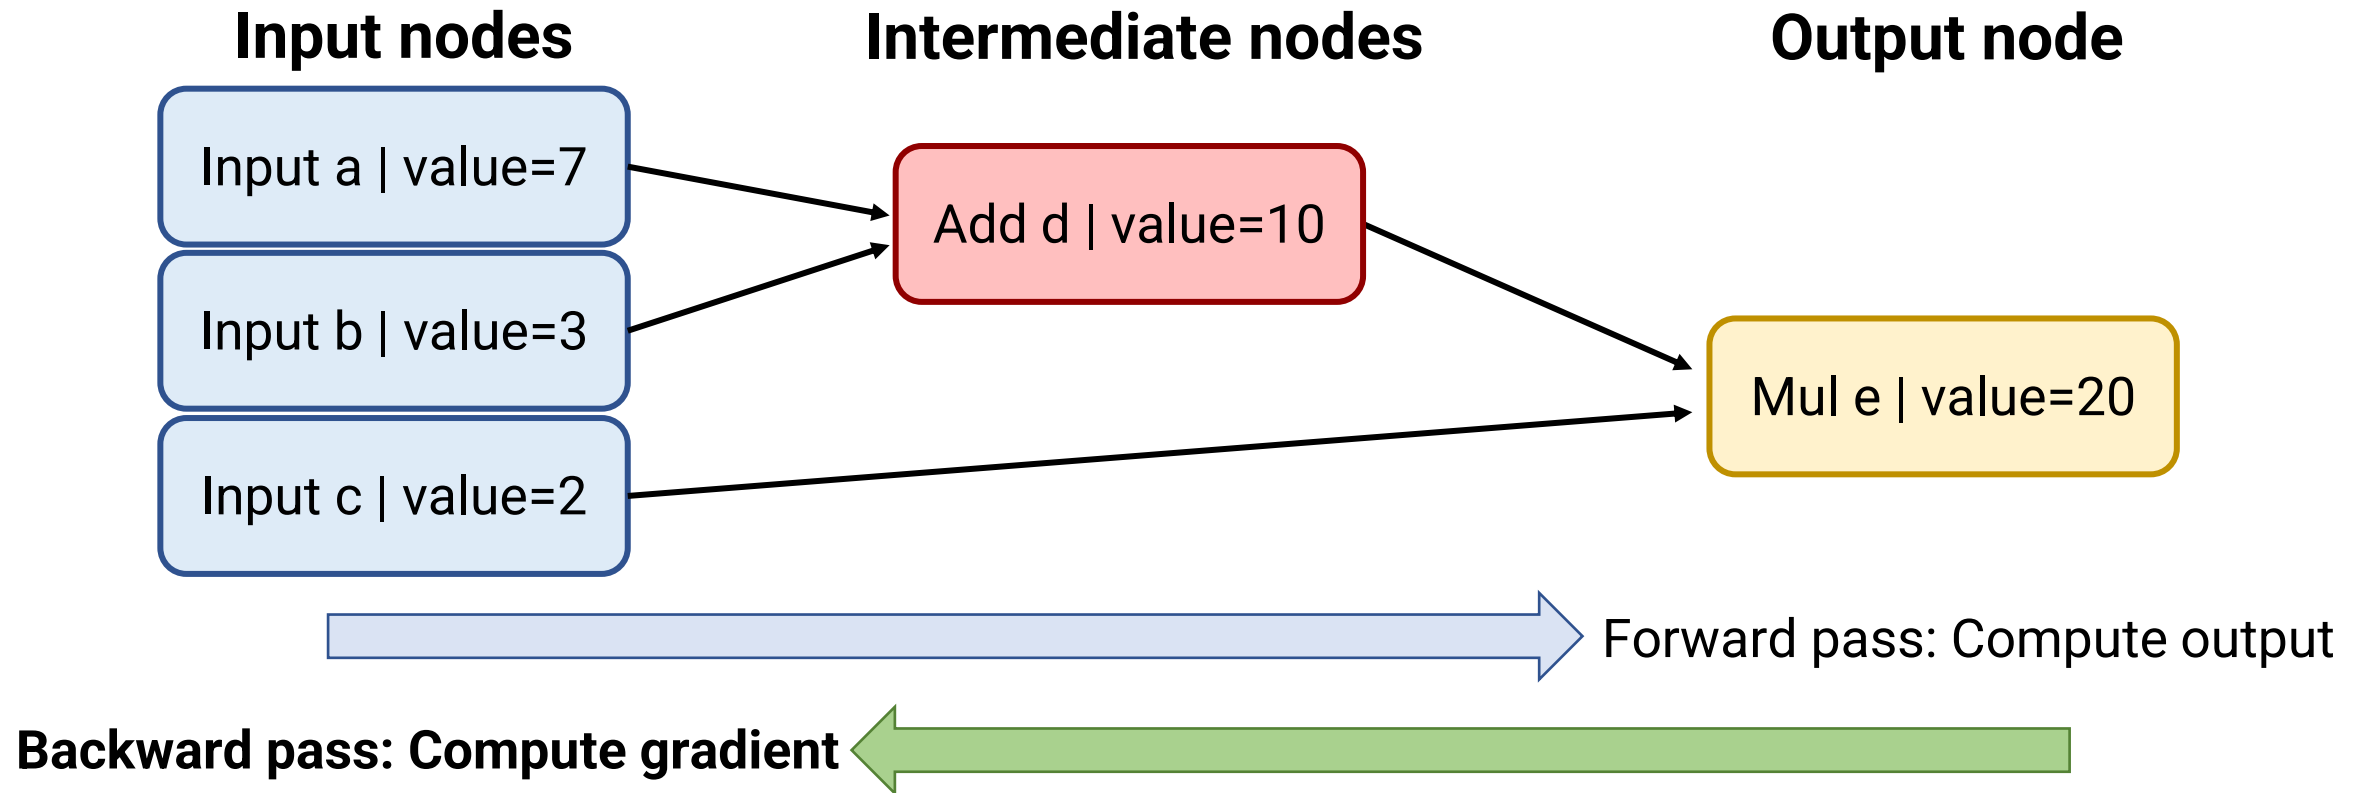
# Today's Plan

---

- The computation graph
- **Backpropagation on trees**
- Backpropagation on DAGs

# Computation Graph


Computation graph for  $(a + b) * c$  when  $a=7, b=3, c=2$





# Backpropagation on trees


---



- For now: assume that the computation graph is a tree
  - Each node is only used in a single computation
  - Root of tree is output
  - Leaves of tree are inputs
- Idea: Recursively compute  $\partial(\text{output})/\partial(\text{node})$  for each node, starting at output


# Backpropagation on trees

---



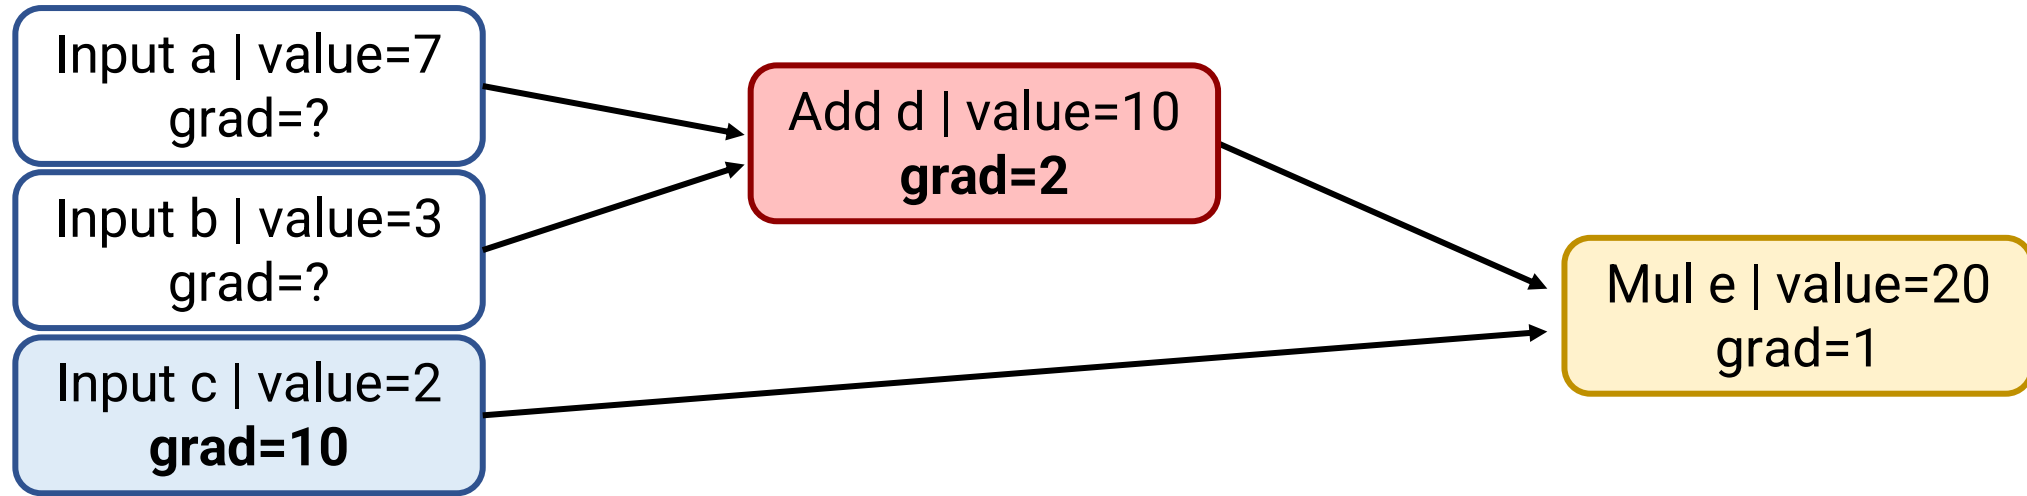
- Goal: Compute gradient  $[\partial e / \partial a, \partial e / \partial b, \partial e / \partial c]$

# Backpropagation on trees



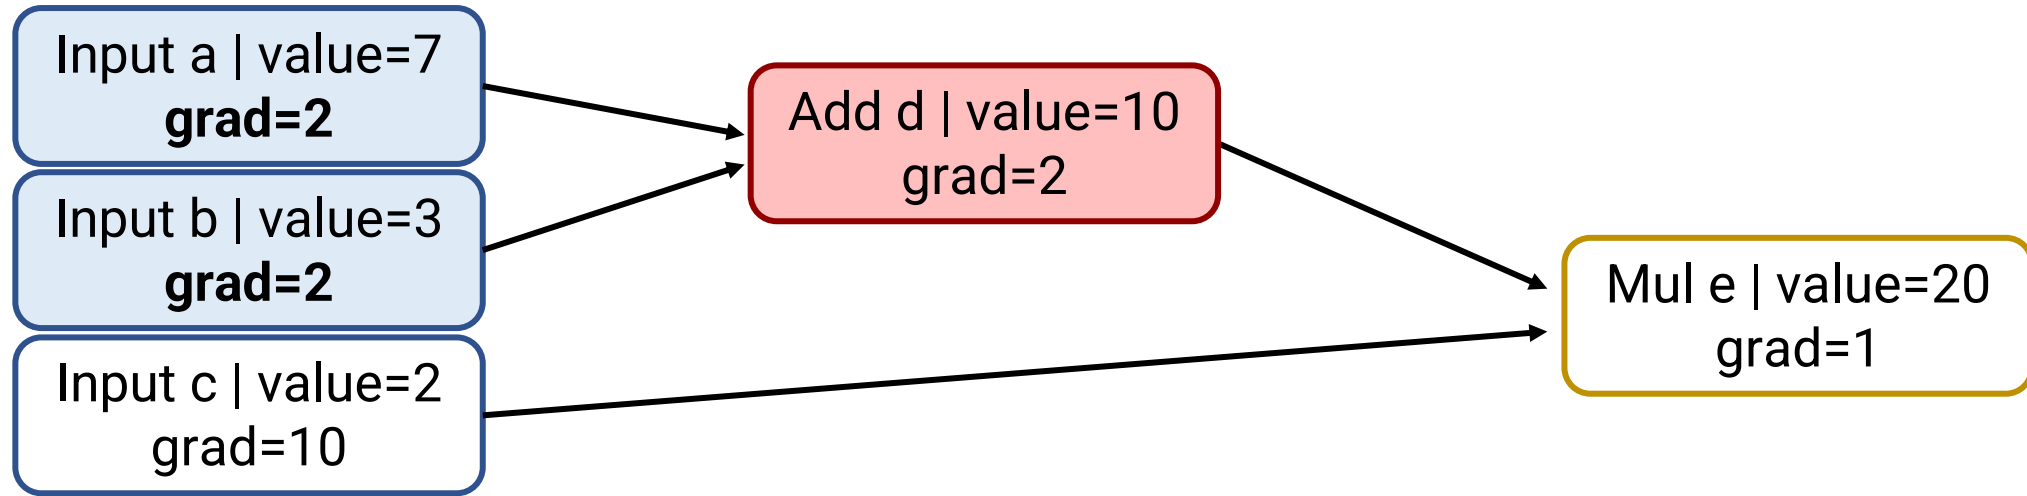
- Goal: Compute gradient  $[\partial e / \partial a, \partial e / \partial b, \partial e / \partial c]$
- Step 1: Base case:  $\partial e / \partial e = 1$

# Backpropagation on trees




- Goal: Compute gradient  $[\partial e / \partial a, \partial e / \partial b, \partial e / \partial c]$
- Step 2: How does **Mul** (node e) “distribute” gradient to its children?
  - $\partial(x*y) / \partial x = y$
  - Chain Rule:  $\partial(out) / \partial x = \partial(out) / \partial(x*y) * \partial(x*y) / \partial x = \partial(out) / \partial(x*y) * y$
  - General rule: Child gets **parent's gradient** \* **value of other child**

# Backpropagation on trees



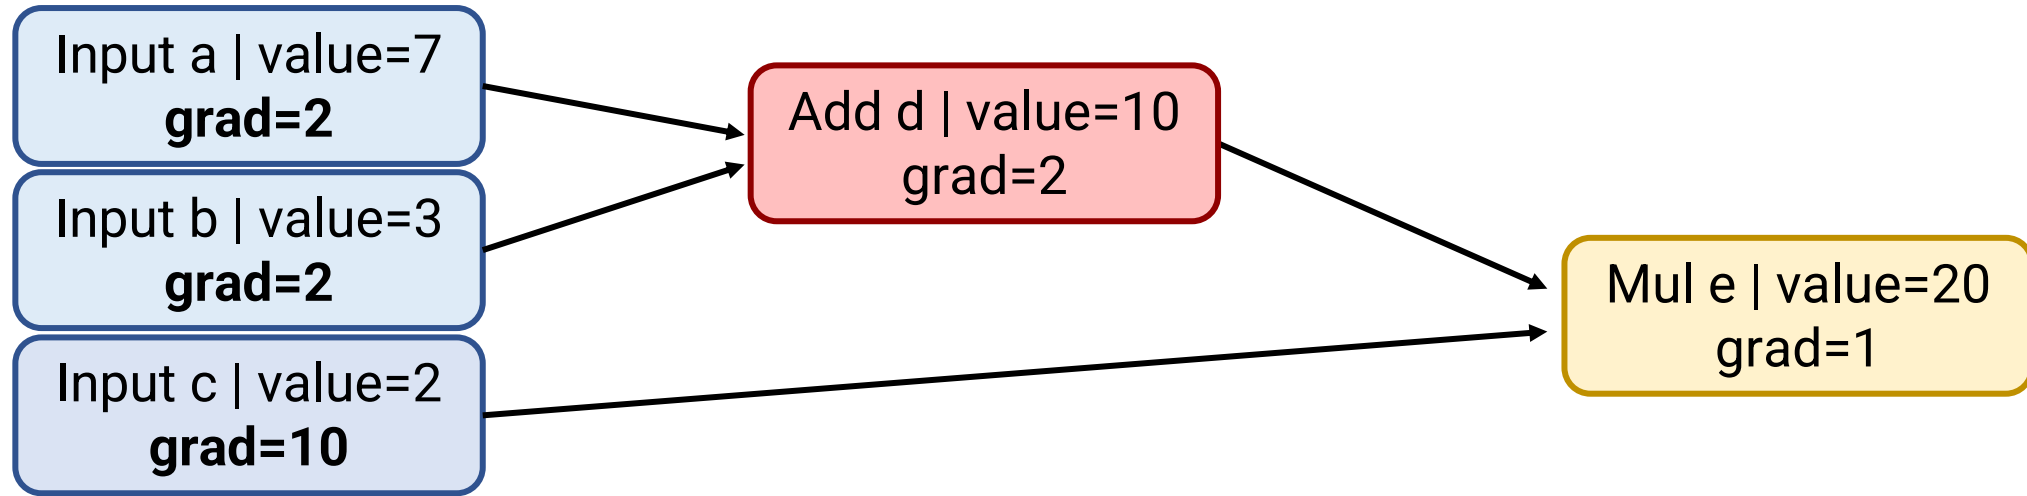
- Goal: Compute gradient  $[\partial e / \partial a, \partial e / \partial b, \partial e / \partial c]$
- Step 3: How does **Add** (node d) “distribute” gradient to its children?
  - $\partial(x+y) / \partial x = 1$
  - Chain Rule:  $\partial(out) / \partial x = \partial(out) / \partial(x+y) * \partial(x+y) / \partial x = \partial(out) / \partial(x+y) * 1$
  - General rule: Child gets **parent’s gradient** \* 1

# Backpropagation on trees



- Goal: Compute gradient  $[\partial e / \partial a, \partial e / \partial b, \partial e / \partial c]$
- Step 4: Leaf nodes
  - Don't need to do anything

# Backpropagation on trees



- Goal: Compute gradient  $[\partial e / \partial a, \partial e / \partial b, \partial e / \partial c]$
- Overall Recipe
  - Do forward pass
  - Start at root and recurse over children
  - Each node knows how to take **gradient of itself with respect to each child**
  - By Chain Rule,  $\text{child.grad} = \text{parent.grad} * \partial(\text{parent}) / \partial(\text{child})$

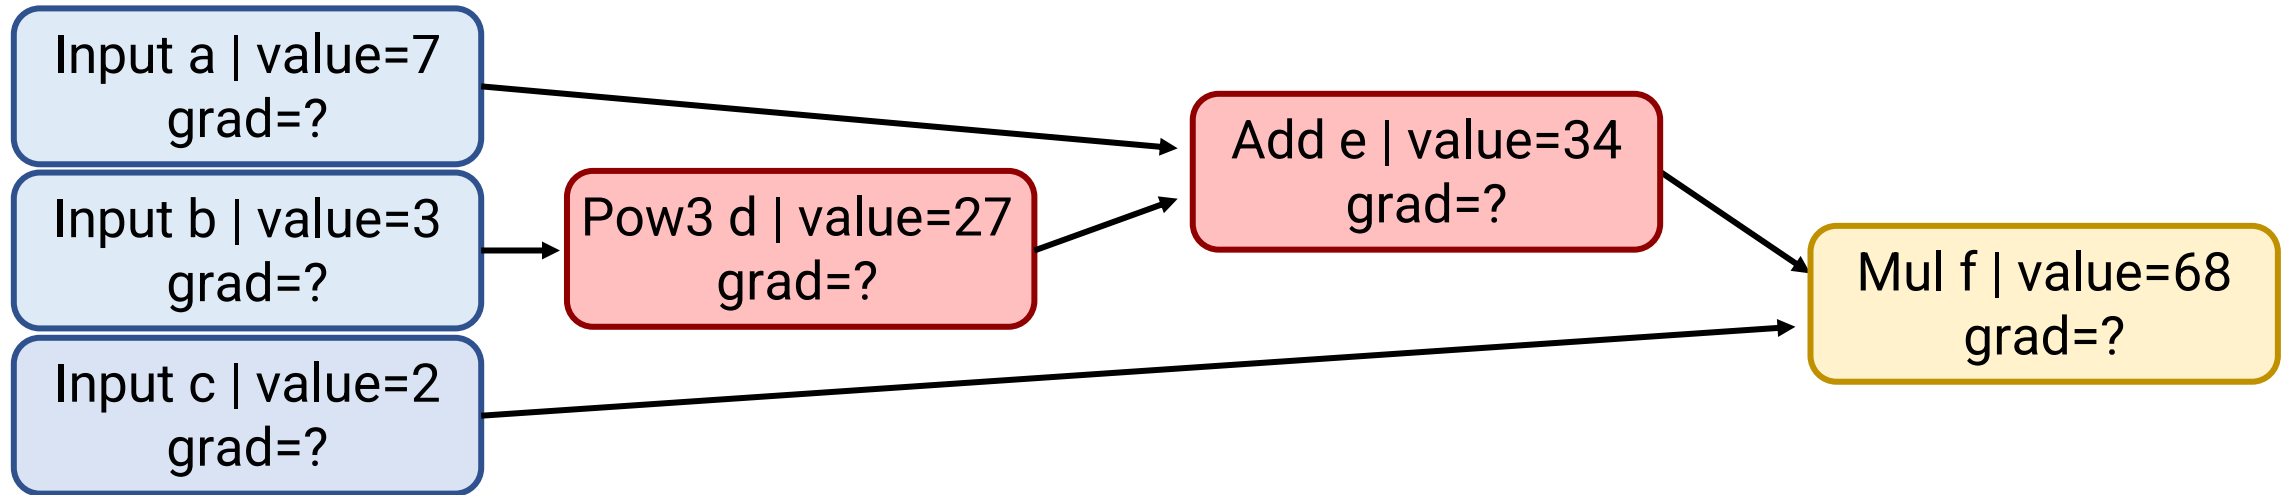
# Let's implement!





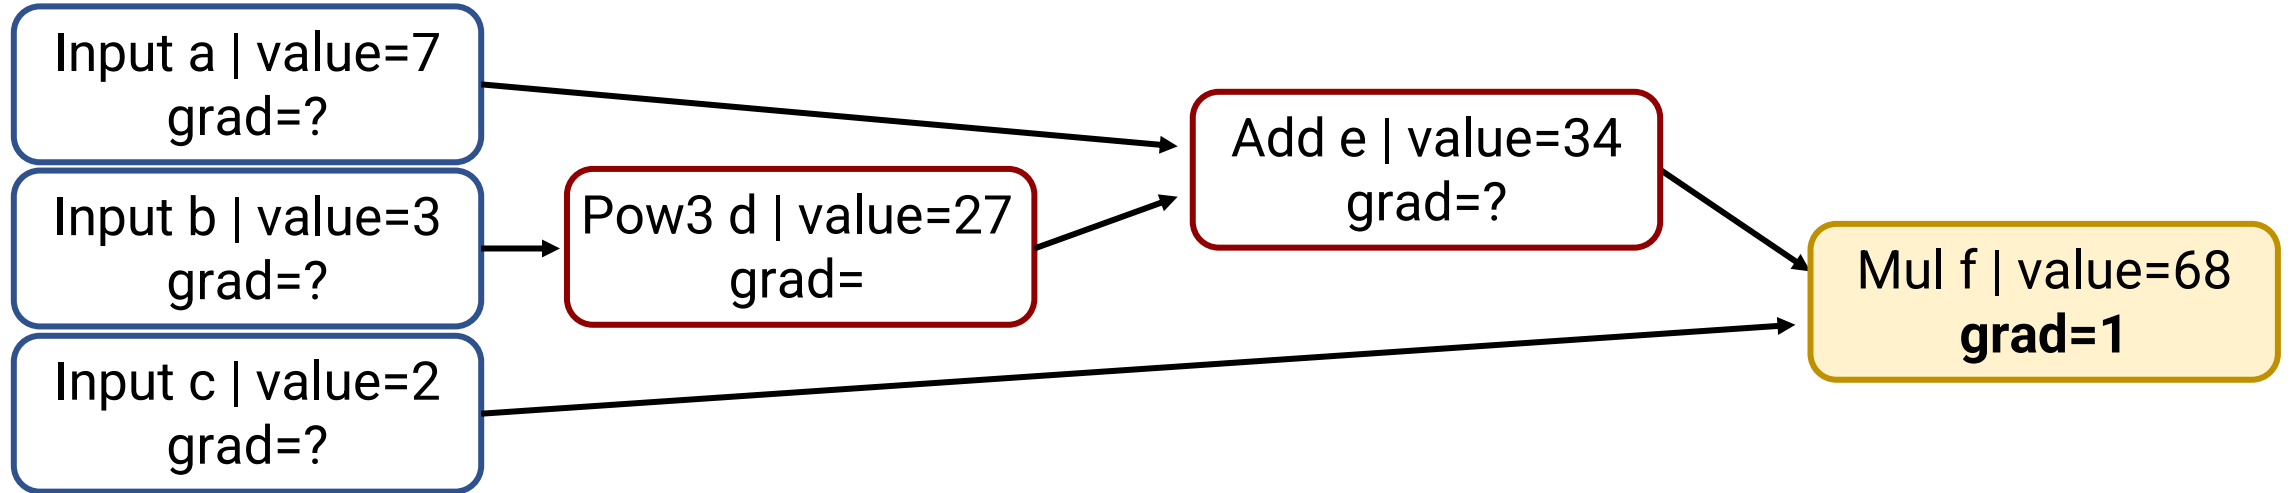
# PowerNode

---



New function:  $(a + b^3) * c$  where  $a=7, b=3, c=2$

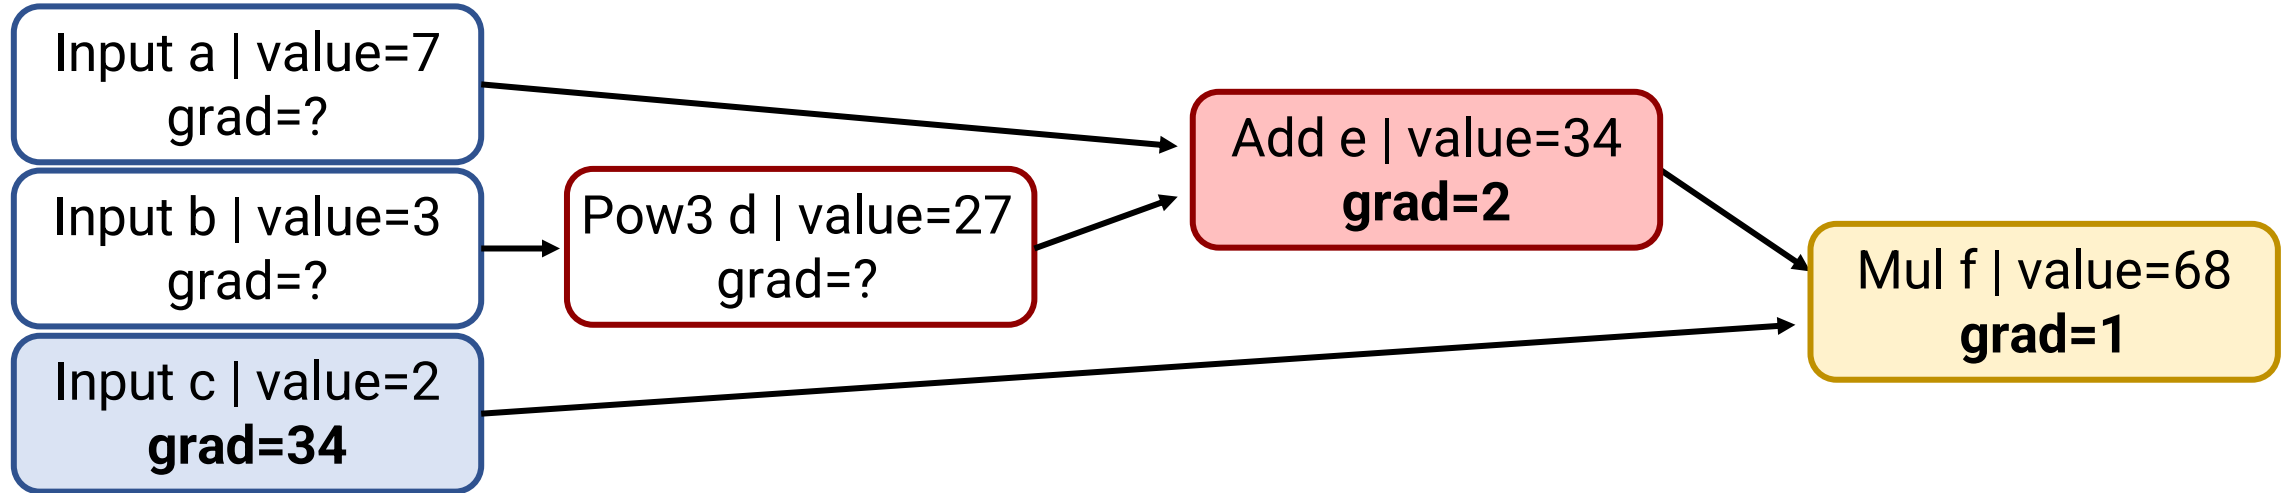
# PowerNode



New function:  $(a + b^3) * c$  where  $a=7, b=3, c=2$

- Goal: Compute gradient  $[\partial f / \partial a, \partial f / \partial b, \partial f / \partial c]$
- Step 1: Base case:  $\partial f / \partial f = 1$

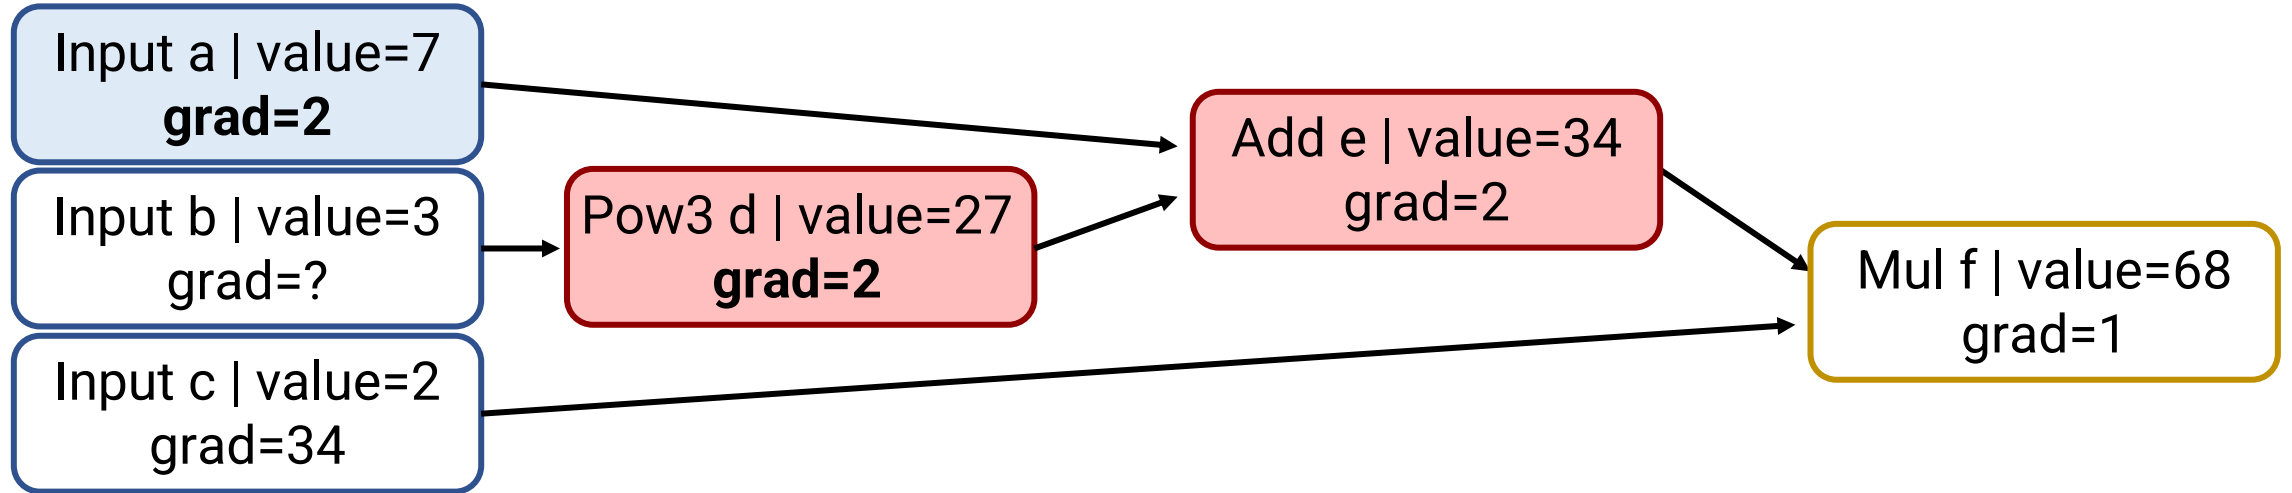
# PowerNode



New function:  $(a + b^3) * c$  where  $a=7, b=3, c=2$

- Goal: Compute gradient  $[\partial f / \partial a, \partial f / \partial b, \partial f / \partial c]$
- Step 2: Distribute **Mul** (node f) gradient to children

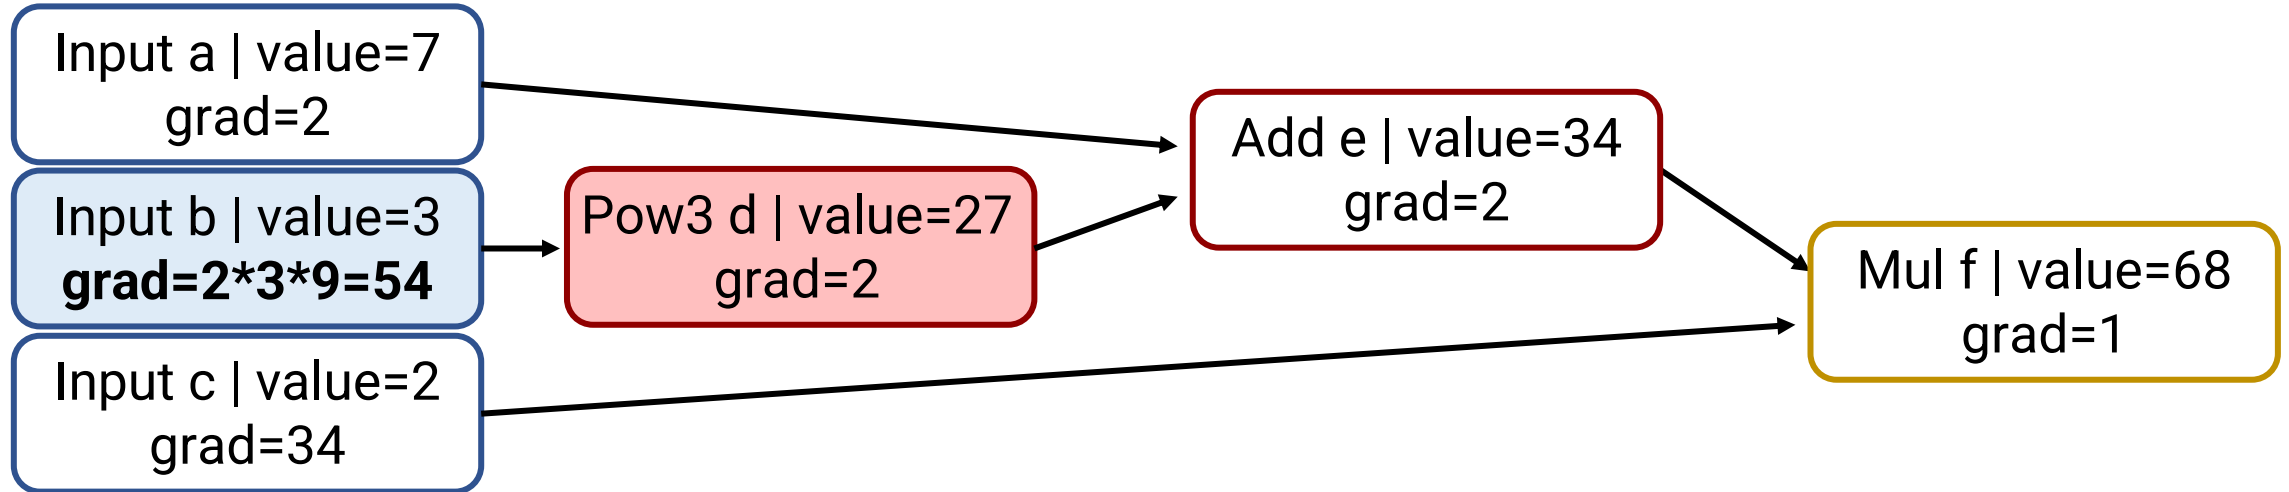
# PowerNode



New function:  $(a + b^3) * c$  where  $a=7, b=3, c=2$

- Goal: Compute gradient  $[\partial f / \partial a, \partial f / \partial b, \partial f / \partial c]$
- Step 3: Distribute **Add** (node e) gradient to children

# PowerNode



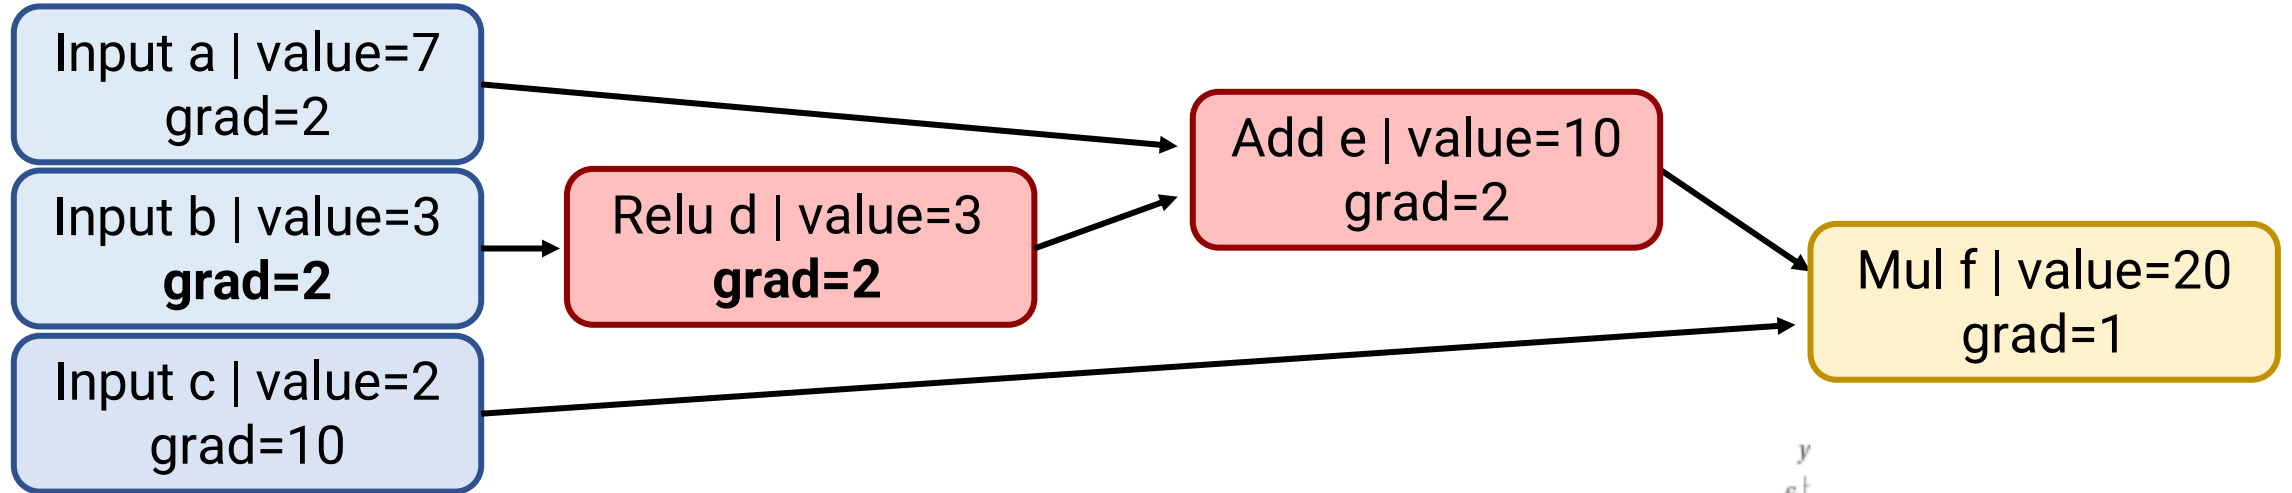
New function:  $(a + b^3) * c$  where  $a=7, b=3, c=2$

- Goal: Compute gradient  $[\partial f / \partial a, \partial f / \partial b, \partial f / \partial c]$
- Step 4: Distribute **Pow3** (node d) gradient to children
  - $\partial(x^p) / \partial x = p * x^{p-1}$
  - By Chain Rule: Child gets **parent's gradient** \*  $p * child^{p-1}$

# Let's implement!




# ReluNode



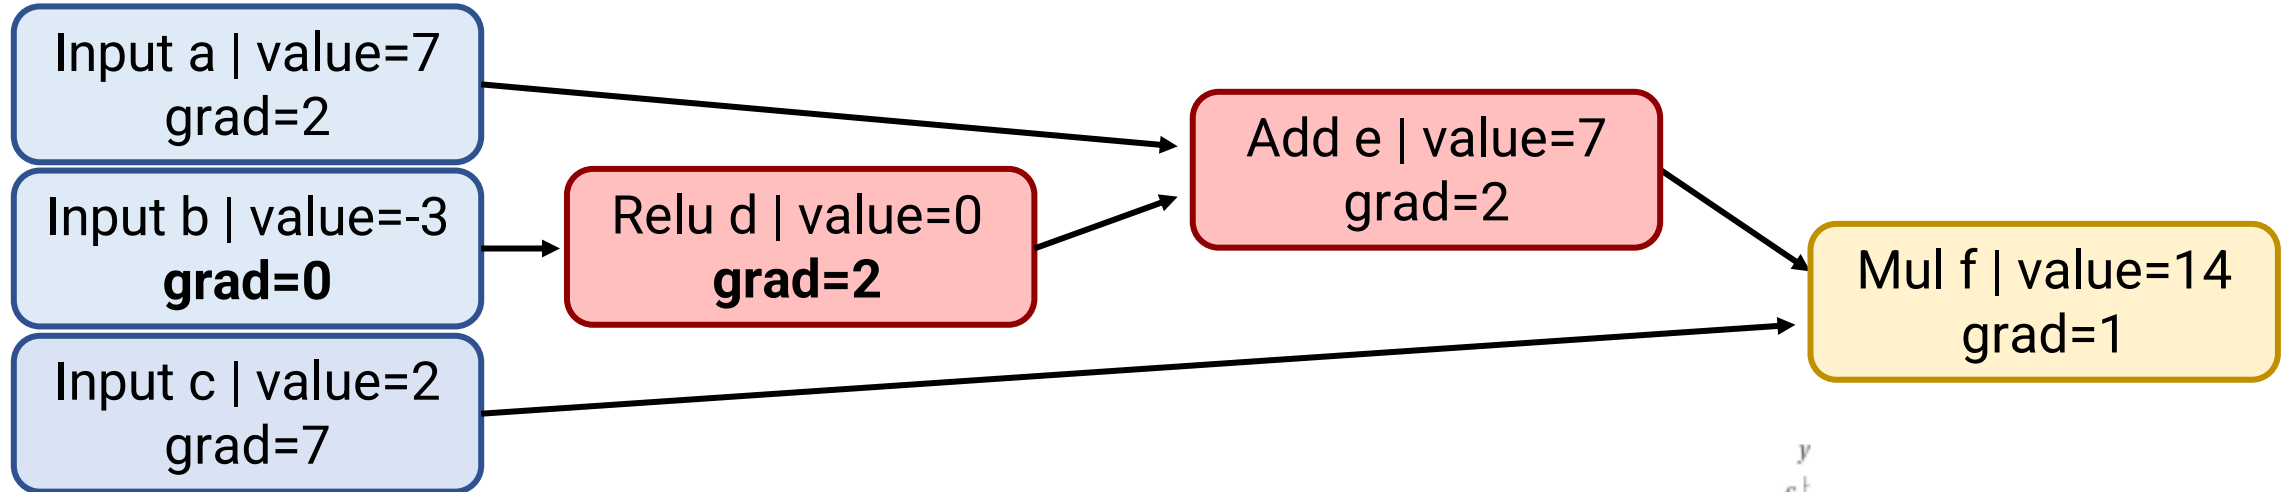
New function:  $(a + \text{Relu}(b)) * c$  where  $a=7, b=3, c=2$

- Steps 1-3 are the same
- Step 4: Relu
  - $\partial(\text{Relu}(x))/\partial x = 1$  if  $x > 0$ ,  $0$  if  $x \leq 0$
  - If child  $> 0$ , child.grad = parent.grad \* 1
  - If child  $\leq 0$ , child.grad = 0




$$\text{ReLU}(z) = \max(z, 0)$$

# ReluNode



New function:  $(a + \text{Relu}(b)) * c$  where  $a=7$ ,  $b=-3$ ,  $c=2$

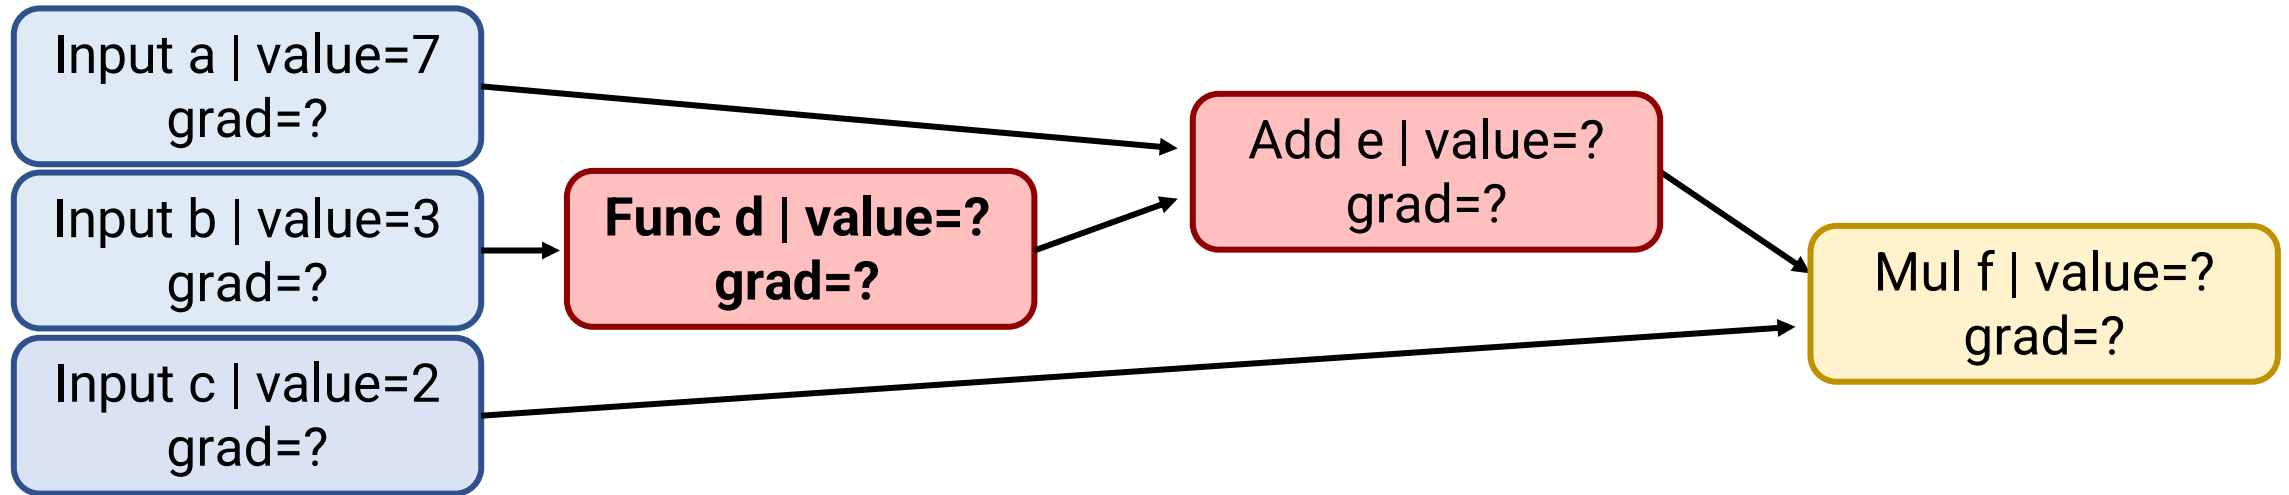
- Steps 1-3 are the same
- Step 4: Relu
  - $\frac{\partial(\text{Relu}(x))}{\partial x} = 1$  if  $x > 0$ ,  $0$  if  $x \leq 0$
  - If child  $> 0$ , child.grad = parent.grad \* 1
  - If child  $\leq 0$ , child.grad = 0



$$\text{ReLU}(z) = \max(z, 0)$$



# Generic Unary Function



New function:  $(a + \text{Func}(b)) * c$  where  $a=7, b=3, c=2$

- Steps 1-3 are the same
- Step 4: Func (generic function)
  - $\text{child.grad} = \text{parent.grad} * \partial(\text{Func}(\text{child}))/\partial \text{child}$

# Announcements

---

- Project proposals due today @ 11:59pm
  - Submit as a group on Gradescope, list all teammates
  - One submission per group
- HW2 released, due Thursday, March 6
- Midterm exam Thursday, March 13
  - In-class, 80 minutes (room TBD)
  - Allowed one double-sided 8.5x11 sheet of notes
  - I highly recommend writing this yourself (good for memory)
- Section Friday: Pytorch (library that does backpropagation)

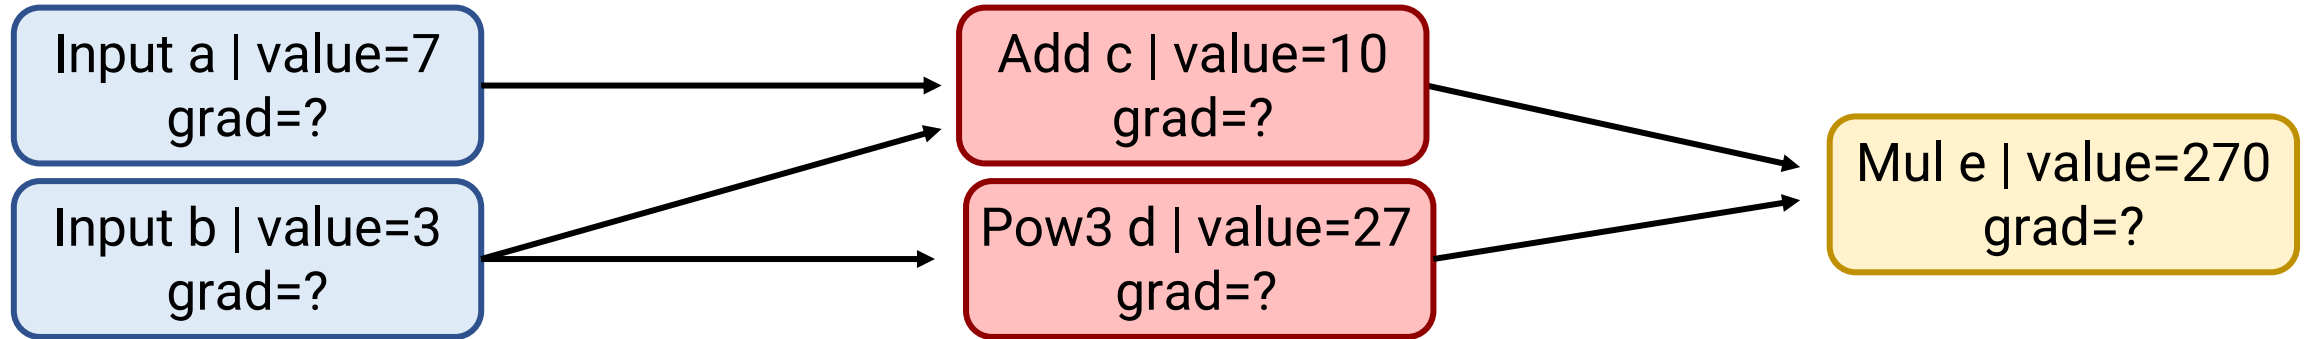
# Today's Plan

---

- The computation graph
- Backpropagation on trees
- **Backpropagation on DAGs**

# DAG Computation Graphs

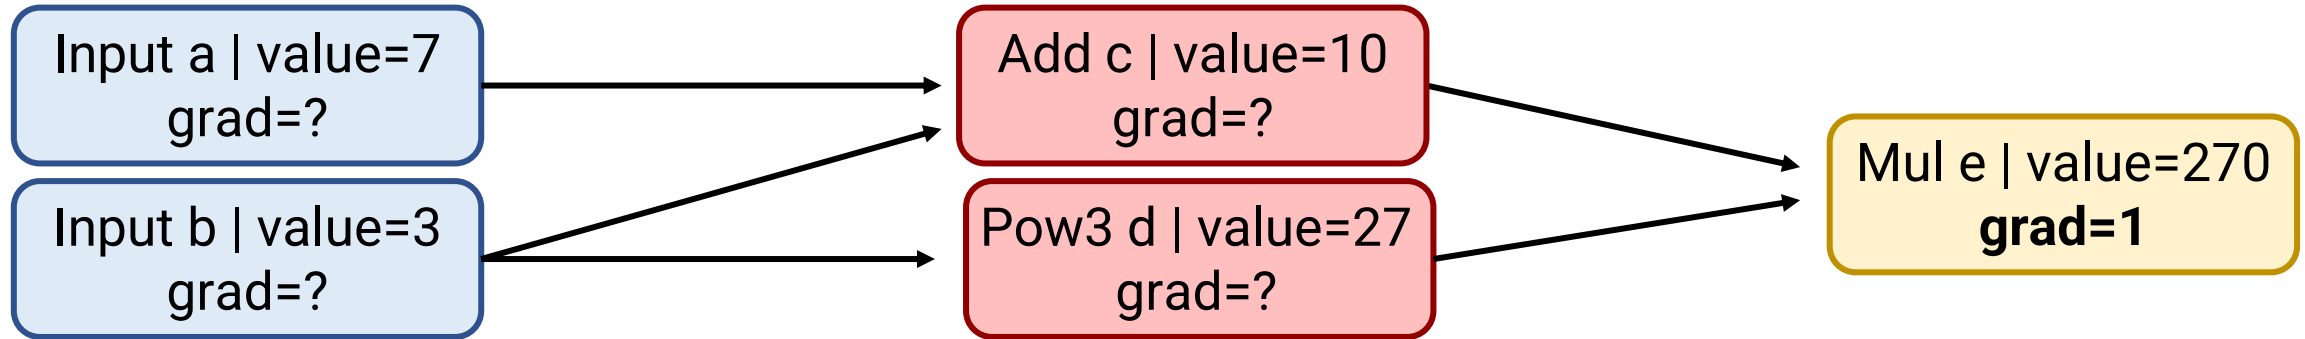
---



New function:  $(a + b) * b^3$  where  $a=7, b=3$

- This is no longer a tree!
- Still a directed acyclic graph
- Let's see why our previous algorithm fails

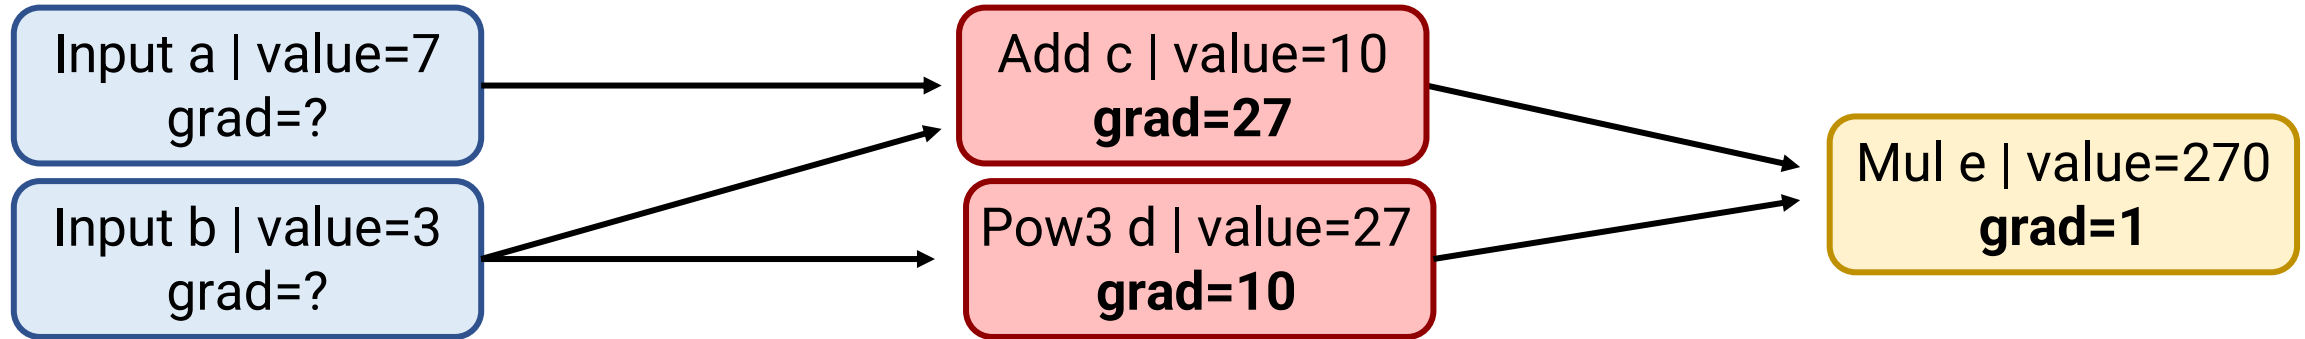
# DAG Computation Graphs



New function:  $(a + b) * b^3$  where  $a=7, b=3$

- Goal: Compute gradient  $[\partial e / \partial a, \partial e / \partial b]$
- Step 1: Base case:  $\partial e / \partial e = 1$

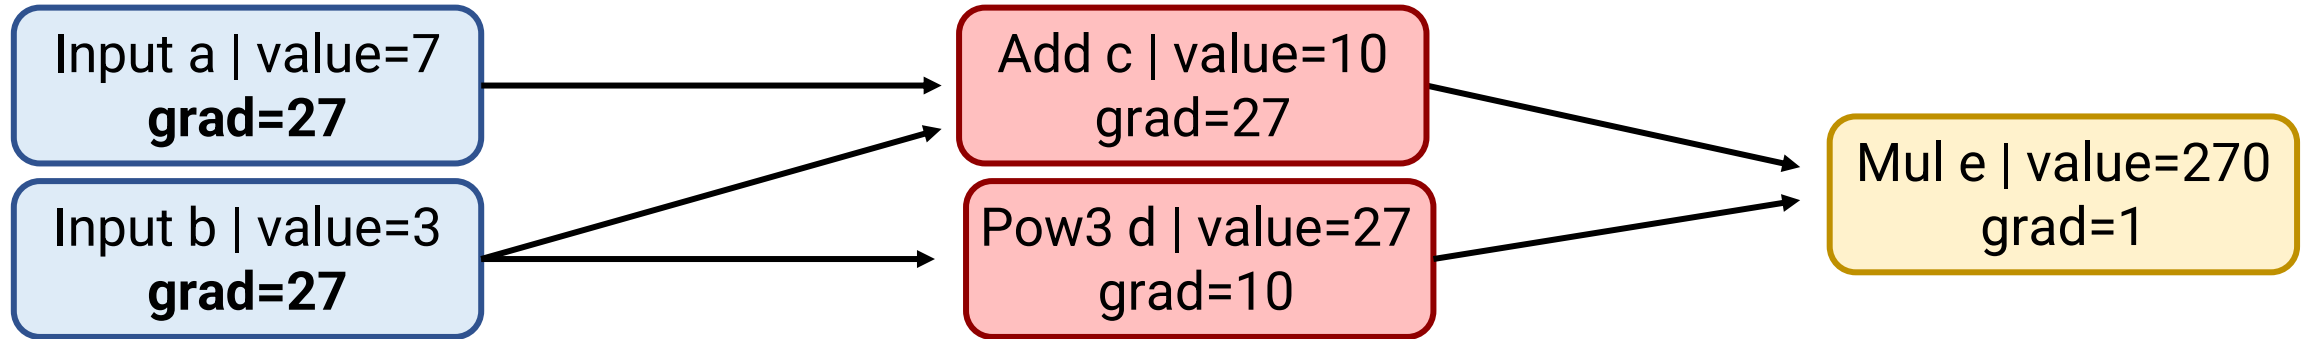
# DAG Computation Graphs



New function:  $(a + b) * b^3$  where  $a=7, b=3$

- Goal: Compute gradient  $[\partial f / \partial a, \partial f / \partial b]$
- Step 2: Distribute **Mul** (node e) gradient to children

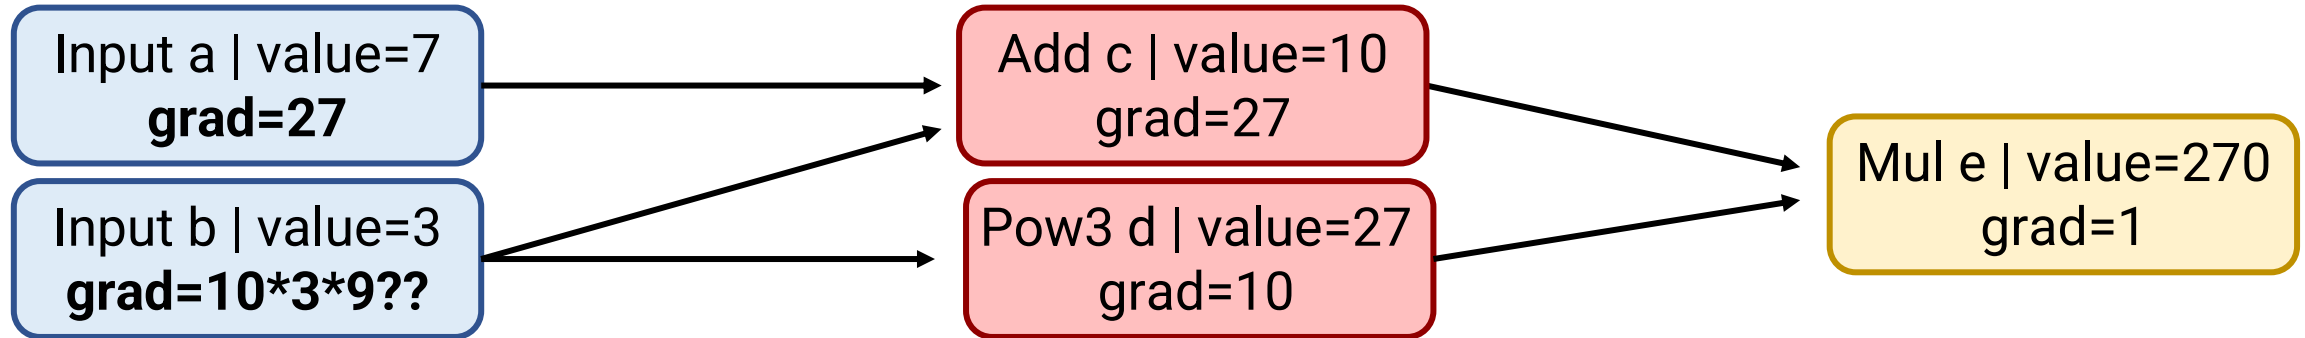
# DAG Computation Graphs



New function:  $(a + b) * b^3$  where  $a=7, b=3$

- Goal: Compute gradient  $[\partial e / \partial a, \partial e / \partial b]$
- Step 3: Distribute **Add** (node c) gradient to children

# DAG Computation Graphs




New function:  $(a + b) * b^3$  where  $a=7, b=3$

- Goal: Compute gradient  $[\partial e / \partial a, \partial e / \partial b]$
- Step 4: Distribute **Pow3** (node d) gradient to child
  - By Chain Rule: Child gets **parent's gradient** \*  $p$  \*  $child^{p-1}$
- Problem: We have overwritten the gradient from b to Add (node c)!



# Multivariate chain rule



- Minimal example


- $u_1$  and  $u_2$  depend on variable  $x$  (i.e.,  $x$  has two parents)
  - $y$  depends on both  $u_1$  and  $u_2$

Dot product of  
 $[\partial y / \partial u_1, \partial y / \partial u_2]$  &  
 $[\partial u_1 / \partial x, \partial u_2 / \partial x]$

- By Multivariate Chain Rule:  $\partial y / \partial x = \partial y / \partial u_1 * \partial u_1 / \partial x + \partial y / \partial u_2 * \partial u_2 / \partial x$ 
  - Changing  $x$  by epsilon changes each parent a bit; Effects add for small epsilon
  - Generalization of multiplying derivatives is matrix multiplication of Jacobians


# Multivariate chain rule

---



- Minimal example
  - $u_1$  and  $u_2$  depend on variable  $x$  (i.e.,  $x$  has two parents)
  - $y$  depends on both  $u_1$  and  $u_2$
- By Multivariate Chain Rule:  $\partial y / \partial x = \partial y / \partial u_1 * \partial u_1 / \partial x + \partial y / \partial u_2 * \partial u_2 / \partial x$
- `node.grad` = **sum over parents** of `parent.grad` \*  $\partial(\text{parent}) / \partial(\text{child})$
- At each parent node, run `child.grad += parent.grad *  $\partial(\text{parent}) / \partial(\text{child})$`

# What order of traversal?




## Current code:

- `y.backward()`
  - `u1.backward()`
    - `w.backward()`
      - `x.backward()`
  - `u2.backward()`
    - `w.backward()`
      - `x.backward()`

- Going recursively double-counts
  - First call to `w.backward()` makes final `x.grad` too large
- Solution: **Topological sort the nodes**
  - Iterate in reverse order, starting from output
  - Ensures that we process each node after all of its parents

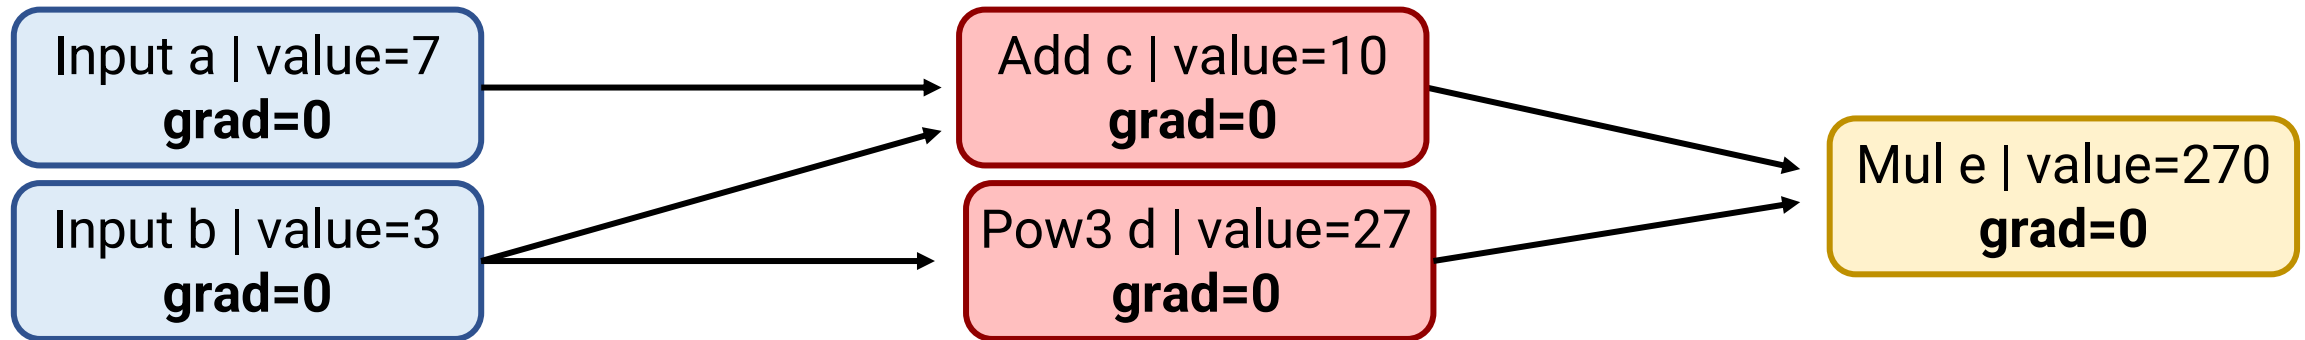
# What order of traversal?



## Better code:

- `topo_order = [x, w, u1, u2, y]`
- Iterate in reverse order:
  - `y.backward()`
  - `u2.backward()`
  - `u1.backward()`
  - `w.backward()`
  - `x.backward()`
- Going recursively double-counts
  - First call to `w.backward()` makes final `x.grad` too large
- Solution: **Topological sort the nodes**
  - Iterate in reverse order, starting from output
  - Ensures that we process each node after all of its parents

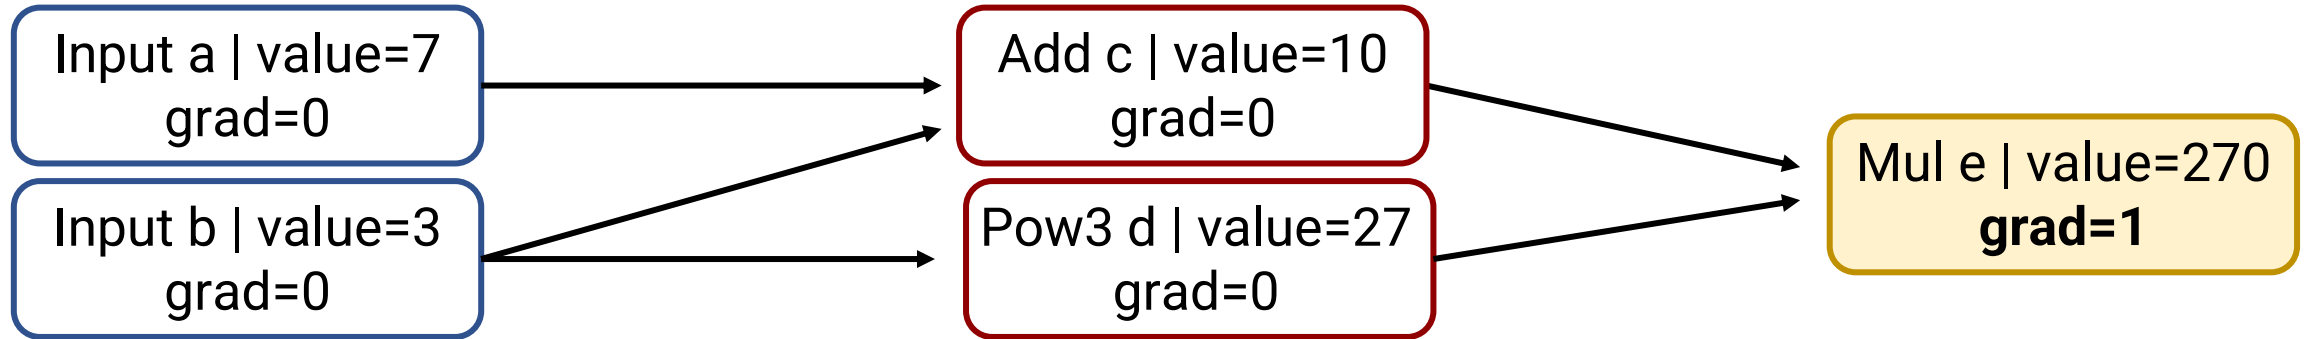
# Backpropagation on DAGs



New function:  $(a + b) * b^3$  where  $a=7, b=3$

- Goal: Compute gradient  $[\partial e / \partial a, \partial e / \partial b]$
- Topological sort:  $[a, b, c, d, e]$
- Step 0: Initialize all gradients to 0


# Backpropagation on DAGs



New function:  $(a + b) * b^3$  where  $a=7, b=3$

- Goal: Compute gradient  $[\partial e / \partial a, \partial e / \partial b]$
- Topological sort:  $[a, b, c, d, e]$
- Step 1: Base case:  $\partial e / \partial e = 1$

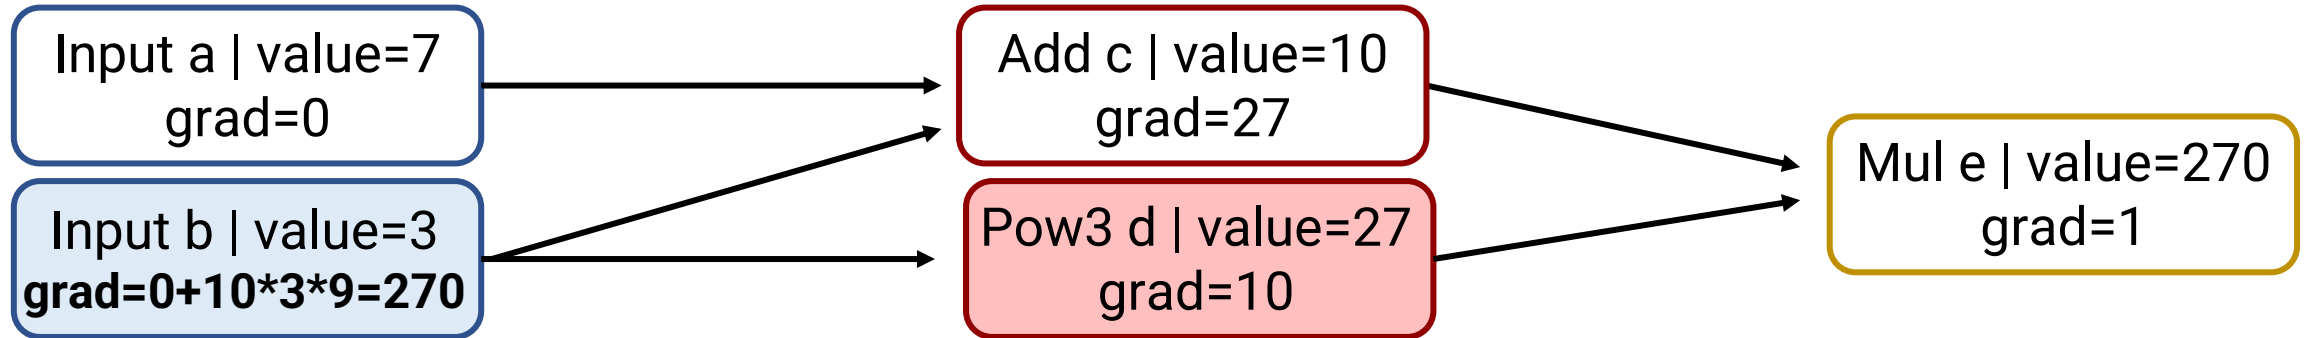
# Backpropagation on DAGs



New function:  $(a + b) * b^3$  where  $a=7, b=3$

- Goal: Compute gradient  $[\partial e / \partial a, \partial e / \partial b]$
- Topological sort, reversed: [e, d, c, b, a]
- Step 2: Propagate **Mul** node e to children

# Backpropagation on DAGs

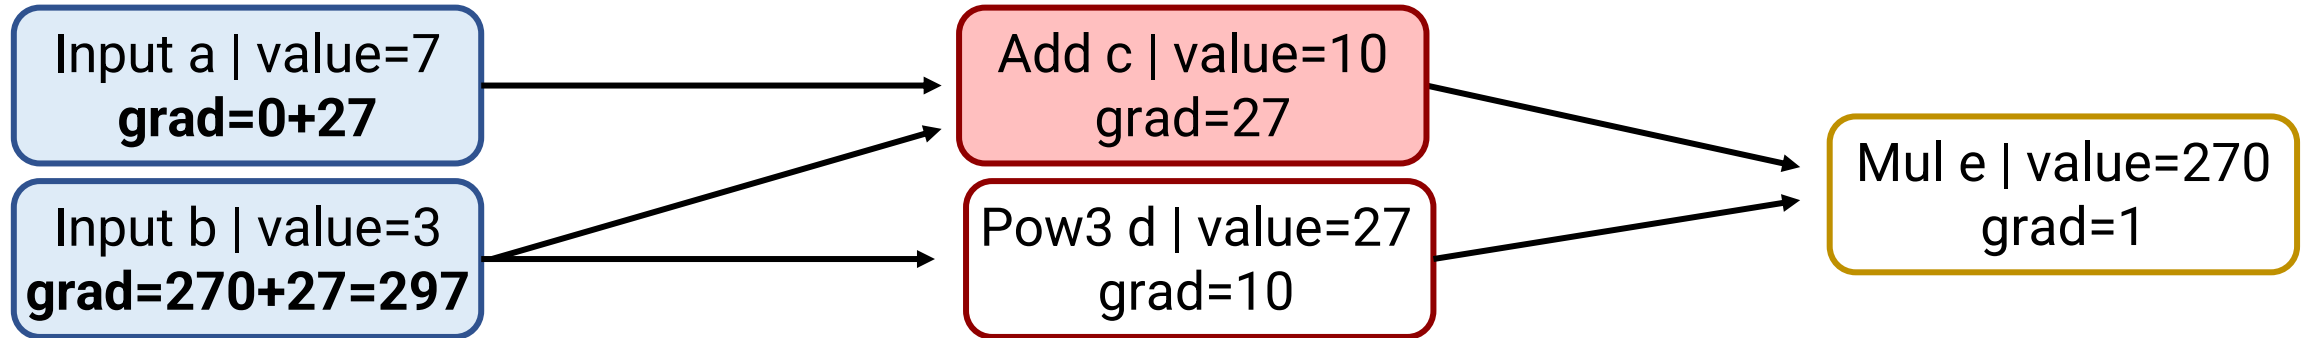


New function:  $(a + b) * b^3$  where  $a=7, b=3$

- Goal: Compute gradient  $[\partial e / \partial a, \partial e / \partial b]$
- Topological sort, reversed:  $[e, d, c, b, a]$
- Step 3: Propagate **Pow3** node d to child



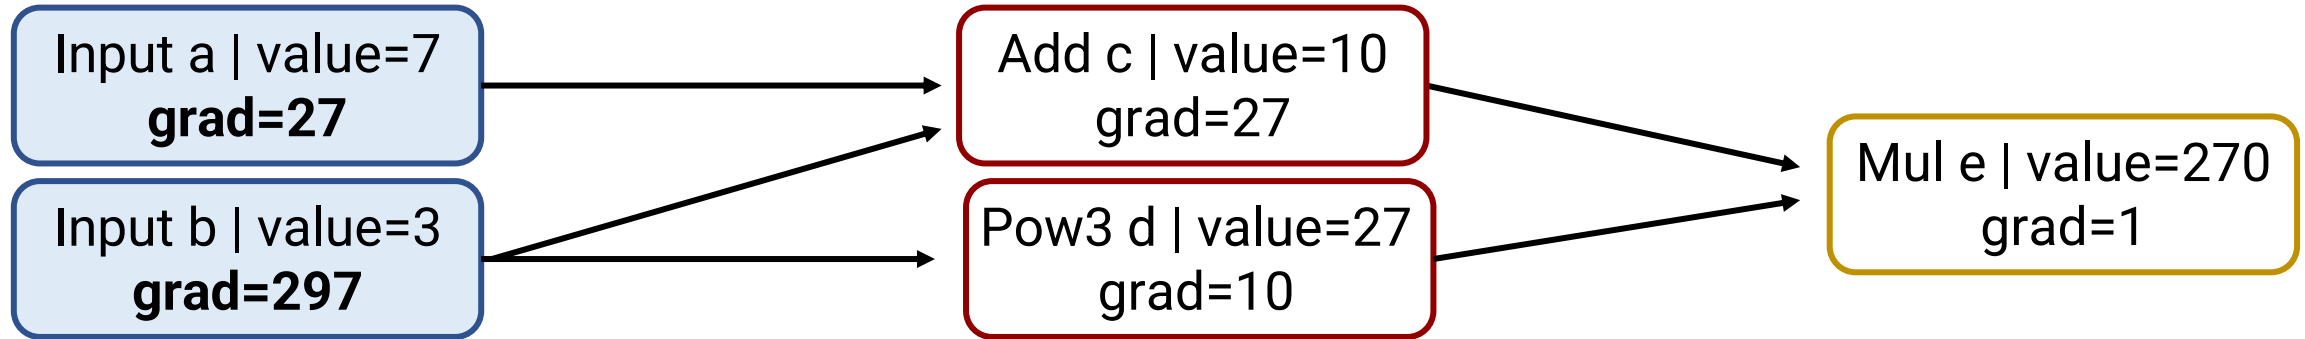
# Backpropagation on DAGs



New function:  $(a + b) * b^3$  where  $a=7, b=3$

- Goal: Compute gradient  $[\partial e / \partial a, \partial e / \partial b]$
- Topological sort, reversed: [e, d, c, b, a]
- Step 4: Propagate **Add** node c to children

# Backpropagation on DAGs



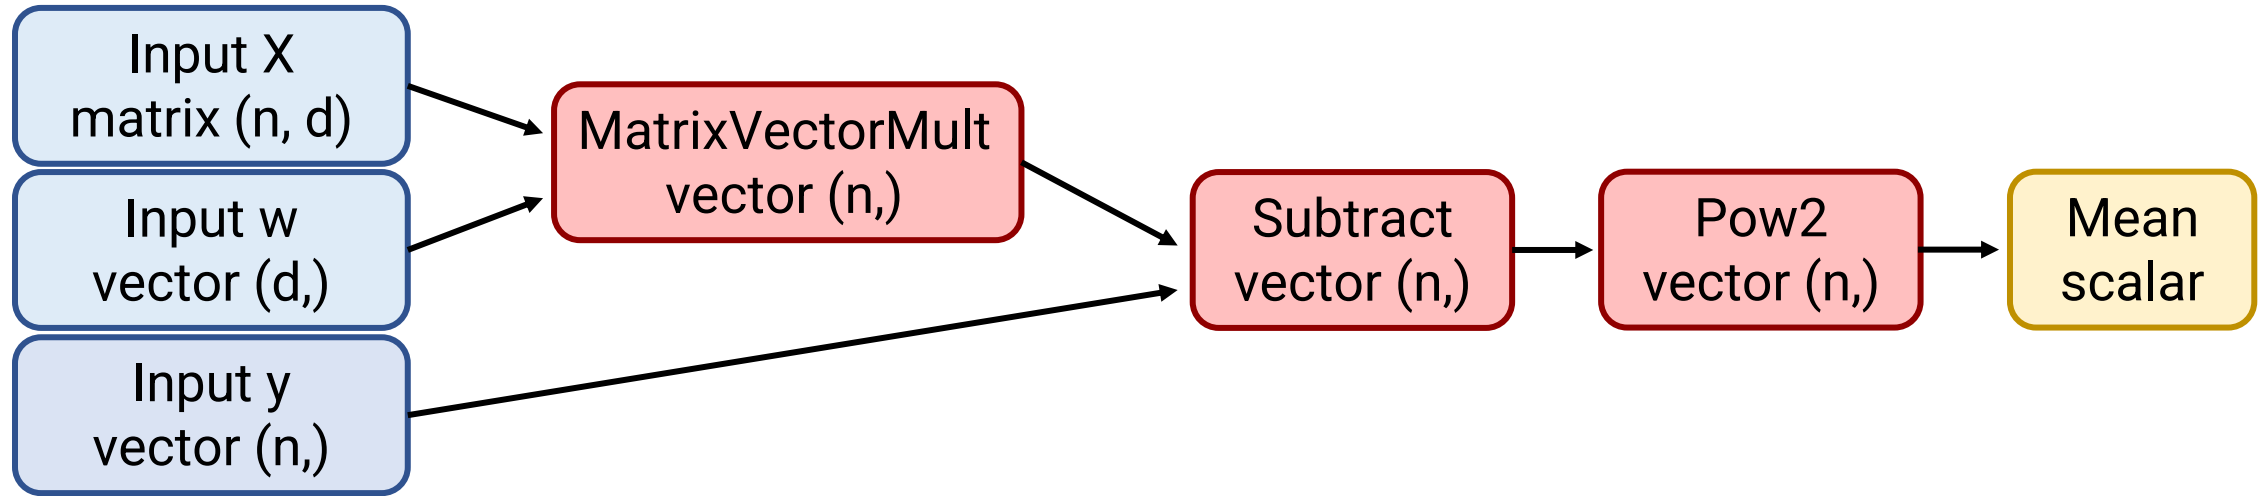
New function:  $(a + b) * b^3$  where  $a=7, b=3$

- Goal: Compute gradient  $[\partial e / \partial a, \partial e / \partial b]$
- Topological sort, reversed:  $[e, d, c, b, a]$
- Step 5, 6: `a.backward()`, `b.backward()` do nothing

# Let's implement!



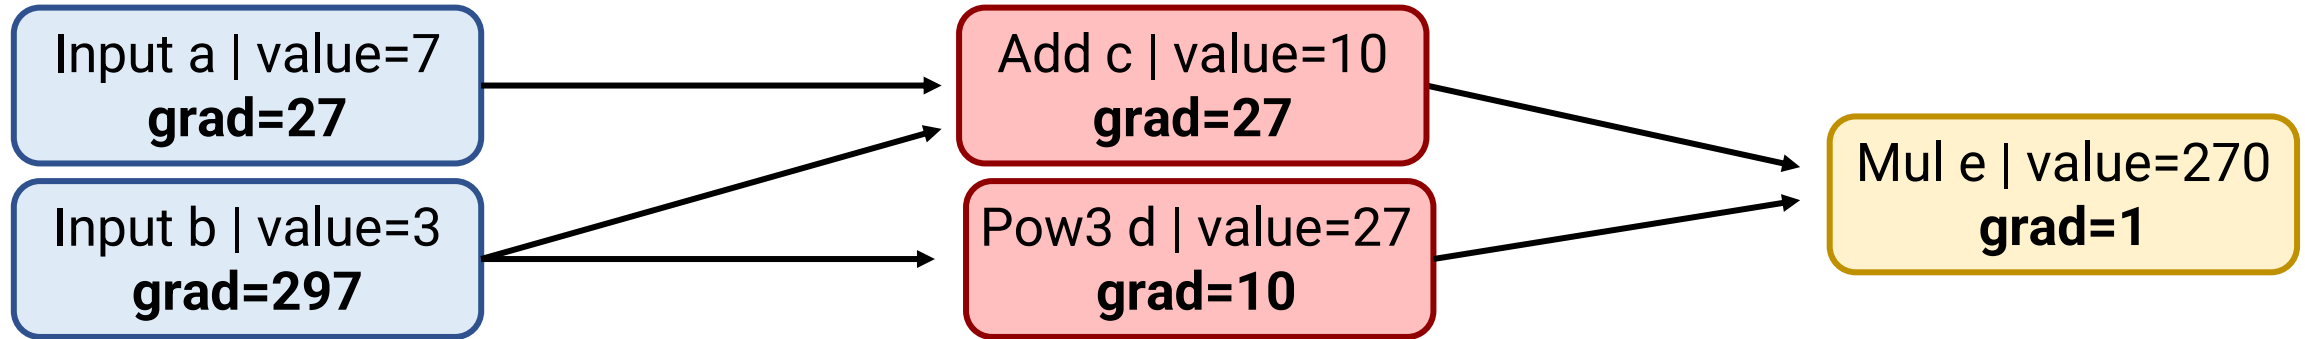
# Backpropagation for vectors and matrices



Computation graph for  $\text{mean}((Xw - y)^2)$ , i.e. Linear Regression

- Basically the same, but each node can be a vector or matrix!
  - Each node.grad stores  $\nabla_{\text{node}} \text{output}$
  - Parents know how to update  $\nabla_{\text{child}} \text{output}$  based on  $\nabla_{\text{parent}} \text{output}$

# Conclusion



- Backpropagation computes gradient of output with respect to all nodes in computation graph
  - Forward pass: Compute values of all nodes
  - Backward pass: Iterate through nodes in reverse order,  
At each parent node, run  $\text{child.grad} += \text{parent.grad} * \frac{\partial(\text{parent})}{\partial(\text{child})}$
- Big picture: Makes it easy to run gradient descent on arbitrary computation graphs
  - Easy to try new architectures for neural networks