

Deep Learning for Images: Convolutional Neural Networks

Robin Jia

USC CSCI 467, Fall 2023


September 28, 2023

Image Classification

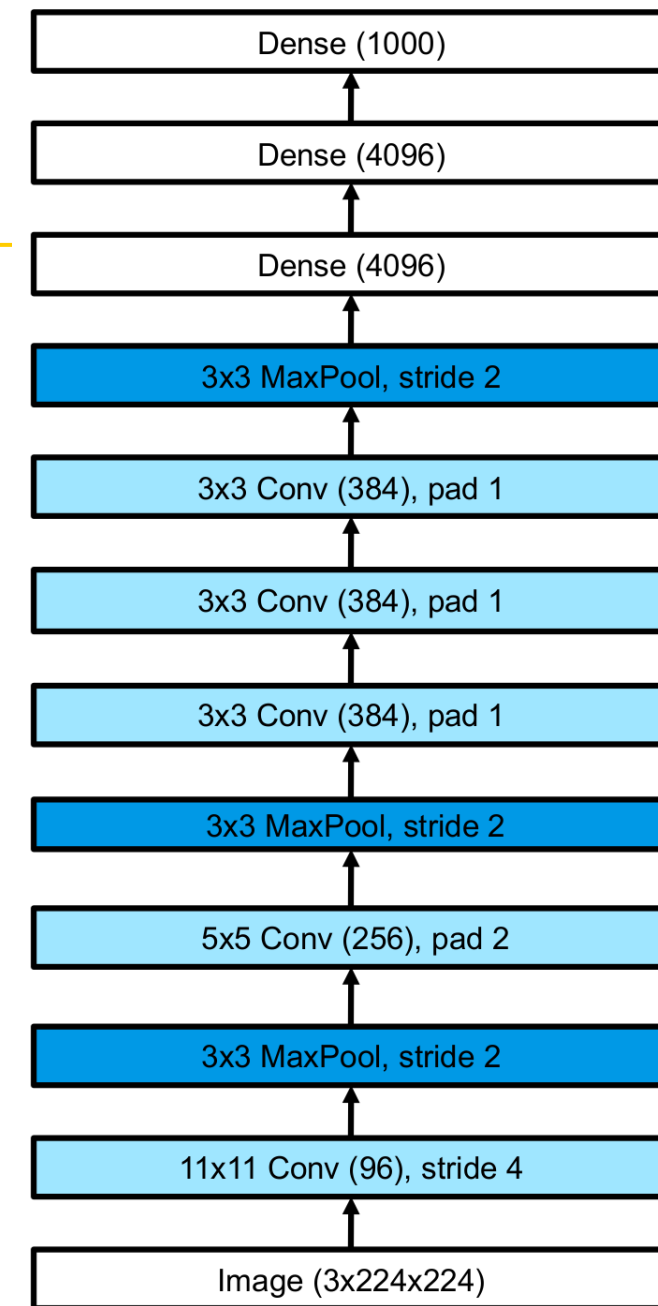


- ImageNet dataset: 14 million images, 1000 labels
- **CNNs do very well at these tasks!**

Progress on ImageNet




- 2012: AlexNet wins ImageNet challenge, marks start of deep learning era **(and is a convolutional neural network)**
- 2016: Machine learning surpasses human accuracy



Outline

- Regularizing Neural Networks
- Intuition for hierarchical features
- Extracting features with convolutions
- Convolutional Neural Networks

Regularization & Neural Networks




- Recall: Neural networks are universal approximators
- This means they are prone to overfitting!
- How to avoid overfitting too badly?

Weight decay (AKA L2 Regularization)

- L2 regularization is a valid strategy!
- Often called “weight decay” when used with neural networks
 - Because every gradient step, you add the update


$$\theta \leftarrow \theta - \eta \cdot \lambda \cdot \theta \quad \text{Weights literally “decay” by factor of } (1 - \eta\lambda)$$

Early stopping




- Idea: Prevent overfitting by stopping training (gradient descent) before you overfit too much
- How it works
 - Every so often during gradient descent, save “checkpoint” of model parameters and evaluate development loss
 - Remember which checkpoint had best development loss
 - If development loss keeps going up, stop training
- Can be used for any model, but especially common for neural networks
 - For linear models, also common to train all the way to convergence

Dropout



(a) Standard Neural Net



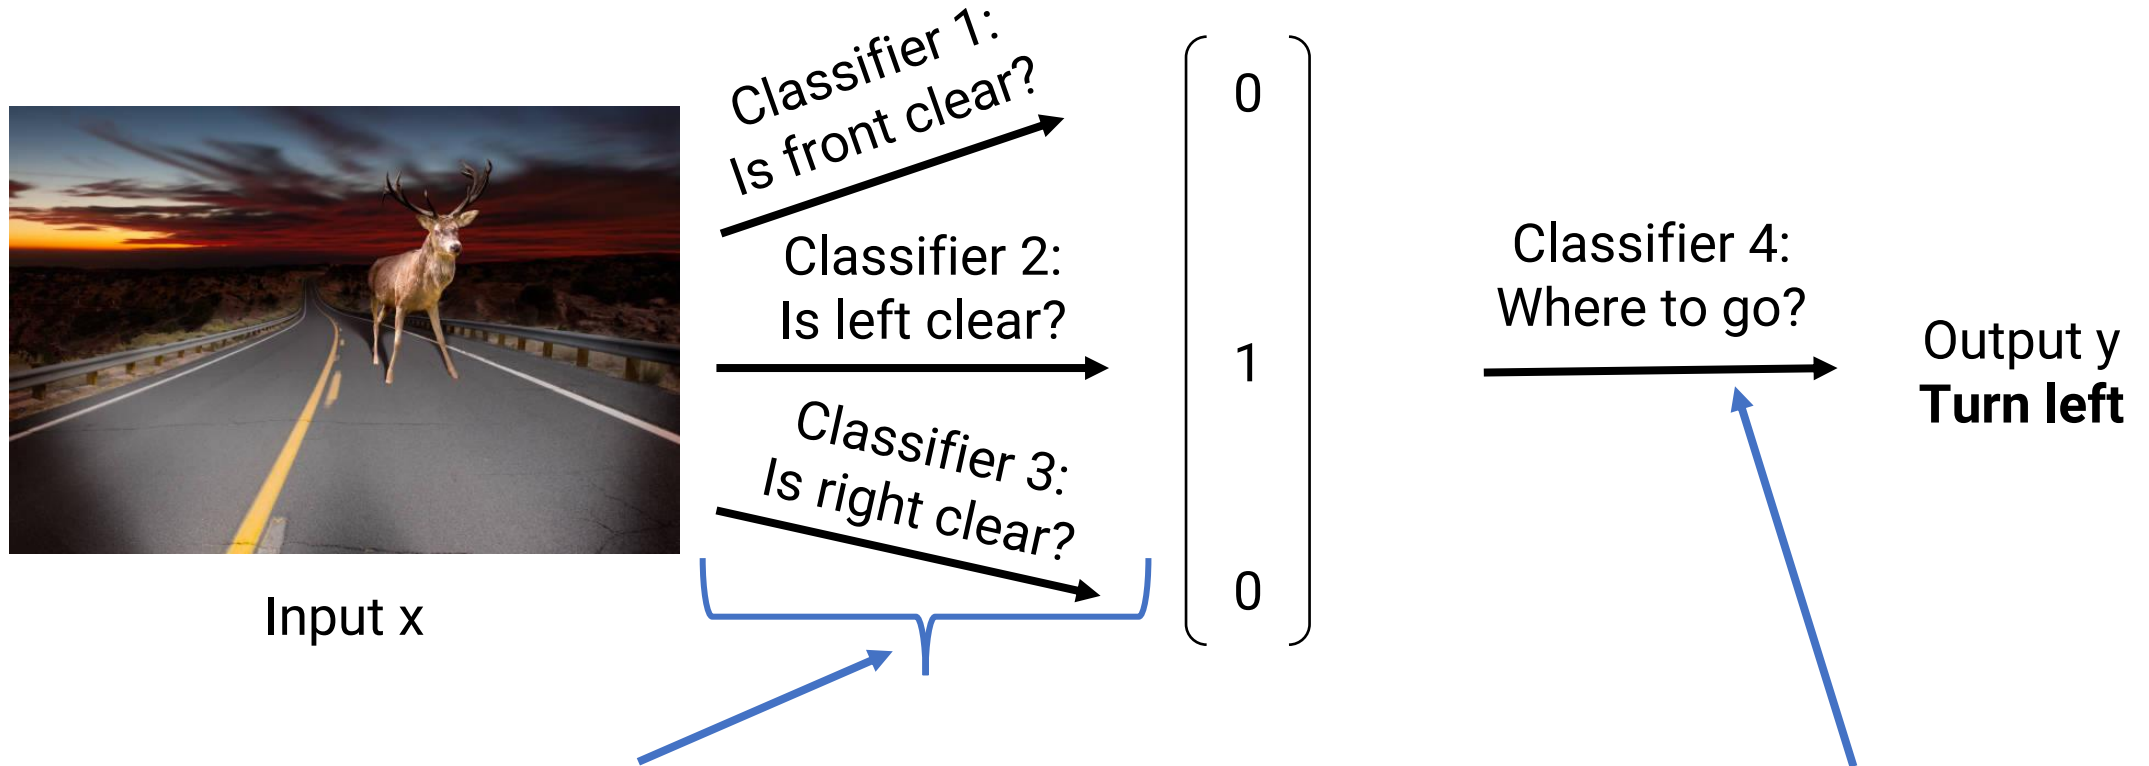
(b) After applying dropout.

- Idea: During training you randomly “drop out” some neurons by setting their value to 0
 - Drop each out with probability p
 - To compensate, scale the other neurons up by $1/p$
 - During testing, don’t do dropout
- Why?
 - Standard intuition is about “co-adaptation” of neurons
 - My personal intuition: Making the problem harder during training is good practice

Outline

- Regularizing Neural Networks
- **Intuition for hierarchical features**
- Extracting features with convolutions
- Convolutional Neural Networks

Review: Neural networks as feature learners



Learn a classifier whose output is a good feature

We don't tell the model what classifier to learn
Model must learn that "is front clear" is a useful concept

Learn to classify based on features
(same as linear model)


A hierarchy of features

- Turn left?




A hierarchy of features

- Turn left?
- **Front is clear?**



A hierarchy of features

- Turn left?
- Front is clear?
- **Is object a moose?**



A hierarchy of features

- Turn left?
- Front is clear?
- Is object a moose?
- **Is this a head?**



A hierarchy of features

- Turn left?
- Front is clear?
- Is object a moose?
- Is this a head?
- **Is this an antler?**




A hierarchy of features

- Turn left?
- Front is clear?
- Is object a moose?
- Is this a head?
- Is this an antler?
- **Is this a line?**



Learning features hierarchically

- Today: **Process images by learning features hierarchically**
- Start with most basic features on smallest patches (e.g., a line)
- Based on those, identify more complex features (e.g., a moose)



Outline

- Regularizing Neural Networks
- Intuition for hierarchical features
- **Extracting features with convolutions**
- Convolutional Neural Networks


A moose detector

- Suppose you have a classifier that can tell if a region has a moose
- How to use it to create a useful feature vector?
- Slide it over each region and check if there's a moose there!




A moose detector

- Suppose you have a classifier that can tell if a region has a moose
- How to use it to create a useful feature vector?
- Slide it over each region and check if there's a moose there!




A moose detector

- Suppose you have a classifier that can tell if a region has a moose
- How to use it to create a useful feature vector?
- Slide it over each region and check if there's a moose there!



A moose detector

- Suppose you have a classifier that can tell if a region has a moose
- How to use it to create a useful feature vector?
- Slide it over each region and check if there's a moose there!
- We just did a convolution!



Learned features $\begin{pmatrix} 0 \\ 0 \\ 1 \\ 0 \\ \dots \end{pmatrix}$

- Moose in far left?
- Moose in center left?
- Moose in center right?
- Moose in far right?

An edge detector

Let's start a little less ambitiously...can we detect a vertical line?

-1	2	-1
-1	2	-1
-1	2	-1

**(Convolutional)
Kernel**

3x3 matrix

0	0	0	0	0	0
0	1	0	0	0	0
0	1	1	1	1	1
0	1	0	0	0	0
0	0	0	0	0	0

Input image
5x6 matrix

Convolve

→
Dot product
kernel &
each image
patch

Output
3x4 matrix

An edge detector

Let's start a little less ambitiously...can we detect a vertical line?

-1	2	-1
-1	2	-1
-1	2	-1

**(Convolutional)
Kernel**

3x3 matrix

0	0	0	0	0	0
0	1	0	0	0	0
0	1	1	1	1	1
0	1	0	0	0	0
0	0	0	0	0	0

Input image
5x6 matrix

Convolve

→
Dot product
kernel &
each image
patch

3			

Output
3x4 matrix

An edge detector

Let's start a little less ambitiously...can we detect a vertical line?

-1	2	-1
-1	2	-1
-1	2	-1

**(Convolutional)
Kernel**

3x3 matrix

0	0	0	0	0	0
0	1	0	0	0	0
0	1	1	1	1	1
0	1	0	0	0	0
0	0	0	0	0	0

Input image
5x6 matrix

Convolve

→
Dot product
kernel &
each image
patch

3	-1		

Output
3x4 matrix

An edge detector

Let's start a little less ambitiously...can we detect a vertical line?

-1	2	-1
-1	2	-1
-1	2	-1

**(Convolutional)
Kernel**

3x3 matrix

0	0	0	0	0	0
0	1	0	0	0	0
0	1	1	1	1	1
0	1	0	0	0	0
0	0	0	0	0	0

Input image
5x6 matrix

Convolve

→
Dot product
kernel &
each image
patch

3	-1	0	

Output
3x4 matrix

An edge detector

Let's start a little less ambitiously...can we detect a vertical line?

-1	2	-1
-1	2	-1
-1	2	-1

**(Convolutional)
Kernel**

3x3 matrix

0	0	0	0	0	0
0	1	0	0	0	0
0	1	1	1	1	1
0	1	0	0	0	0
0	0	0	0	0	0

Input image
5x6 matrix

Convolve

→
Dot product
kernel &
each image
patch

3	-1	0	0

Output
3x4 matrix

An edge detector

Let's start a little less ambitiously...can we detect a vertical line?

-1	2	-1
-1	2	-1
-1	2	-1

**(Convolutional)
Kernel**

3x3 matrix

0	0	0	0	0	0
0	1	0	0	0	0
0	1	1	1	1	1
0	1	0	0	0	0
0	0	0	0	0	0

Input image
5x6 matrix

Convolve

→
Dot product
kernel &
each image
patch

3	-1	0	0
5			

Output
3x4 matrix

An edge detector

Let's start a little less ambitiously...can we detect a vertical line?

-1	2	-1
-1	2	-1
-1	2	-1

**(Convolutional)
Kernel**

3x3 matrix

0	0	0	0	0	0
0	1	0	0	0	0
0	1	1	1	1	1
0	1	0	0	0	0
0	0	0	0	0	0

Input image
5x6 matrix

Convolve

→
Dot product
kernel &
each image
patch

3	-1	0	0
5	-2		

Output
3x4 matrix

An edge detector

Let's start a little less ambitiously...can we detect a vertical line?

-1	2	-1
-1	2	-1
-1	2	-1

**(Convolutional)
Kernel**

3x3 matrix

0	0	0	0	0	0
0	1	0	0	0	0
0	1	1	1	1	1
0	1	0	0	0	0
0	0	0	0	0	0

Input image
5x6 matrix

Convolve

→
Dot product
kernel &
each image
patch

3	-1	0	0
5	-2	0	

Output
3x4 matrix

An edge detector

Let's start a little less ambitiously...can we detect a vertical line?

-1	2	-1
-1	2	-1
-1	2	-1

**(Convolutional)
Kernel**

3x3 matrix

0	0	0	0	0	0
0	1	0	0	0	0
0	1	1	1	1	1
0	1	0	0	0	0
0	0	0	0	0	0

Input image
5x6 matrix

Convolve

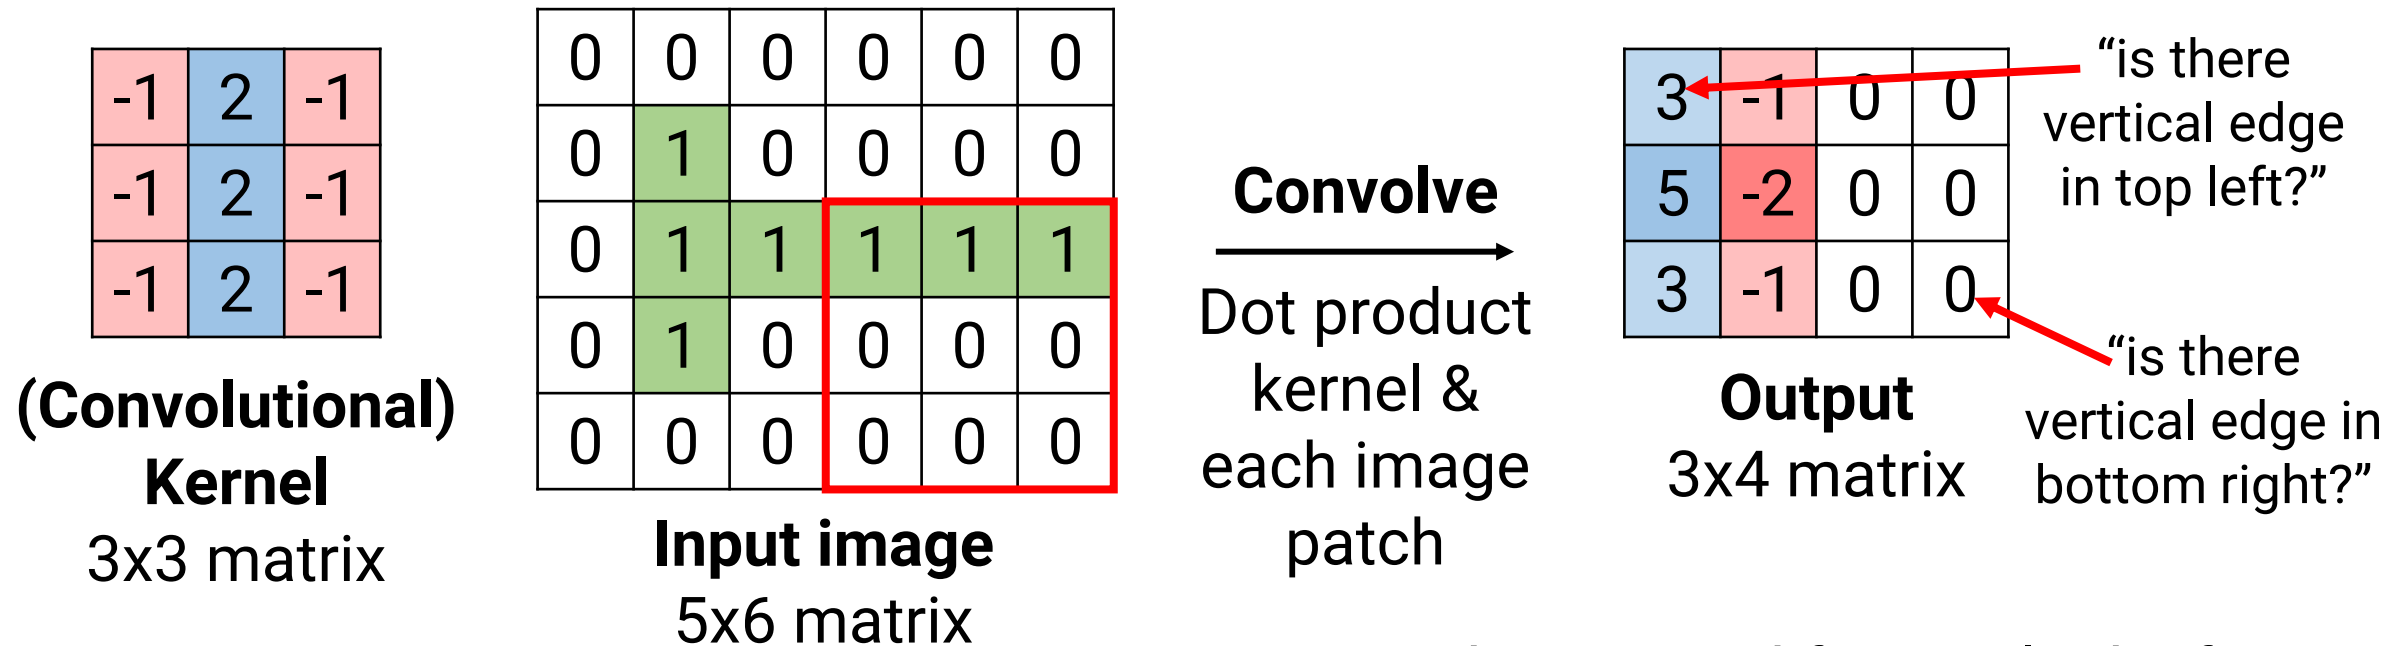
→
Dot product
kernel &
each image
patch

3	-1	0	0
5	-2	0	0

Output
3x4 matrix

An edge detector

Let's start a little less ambitiously...can we detect a vertical line?



Each extracted feature looks for the same thing in different location

Convolutions

-1	2	-1
-1	2	-1
-1	2	-1

Kernel
(K=3)

0	0	0	0	0	0
0	1	0	0	0	0
0	1	1	1	1	1
0	1	0	0	0	0
0	0	0	0	0	0

Input
(5 x 6)
input[1:4,2:5]

3	-1	0	0
5	-2	0	0
3	-1	0	0

Output
(5-3+1 x 6-3+1)
=(3 x 4)

(1, 2)-th
element

- Convolution is an operation that takes in two matrices:
 - Kernel: K x K matrix (e.g., K=3)
 - Input: W x H matrix
- Output: (W-K+1) x (H-K+1) matrix
 - ij-th element of output is dot product of kernel & input[i:i+K,j:j+K]
 - (I'm 0-indexing in these slides)
- Convolutional Layer: Kernel is our weight/parameter, use convolution to extract features
- Note: Convolution is a **linear** operation!


Motivation #1: Local Receptive Fields

- Motivation #1: Each neuron should only look at a small patch of input
- Why? Local textures/shapes are useful
- First understand local patterns, build up to global understanding



Motivation #2: Weight Sharing

- Motivation #2: In each local receptive field, the same types of features are useful
 - Basic: Detecting edges
 - More advanced: Detecting moose
- So, **share the same kernel** (i.e. weights) for all image patches
- Convolutions encode **translation equivariance**
 - If your image gets shifted, convolution outputs just get shifted too



Announcements

- HW2 due next Thursday, October 5
- Midterm exam Tuesday, October 10
 - In-class, 80 minutes, one double-sided 8.5x11 sheet of notes
 - Room assignments (also on Piazza)
 - Last name A-O: LVL 17 (this room)
 - Last name P-Z: THH 116
- Section tomorrow: Pytorch tutorial
 - Important for problem 4 of homework!

Convolutional vs. Fully Connected Layers

-1	2	-1
-1	2	-1
-1	2	-1

Kernel
(size 9)

0	0	0	0	0	0
0	1	0	0	0	0
0	1	1	1	1	1
0	1	0	0	0	0
0	0	0	0	0	0

Input
(size 30)


3	-1	0	0
5	-2	0	0
3	-1	0	0

Output
(size 12)

- Let's count parameters needed
 - Convolutional layer with $K=3$
 - Kernel = $3 \times 3 = 9$ parameters
 - Add a bias term = **10 parameters**
 - Fully connected layer with 30-dim input, 12-dim output needs
 - W : $30 \times 12 = 360$ parameters
 - b : 12 parameters
 - Total: **372 parameters**
- Fewer parameters = need less data to learn useful features
- FC would have to learn to detect the same feature (e.g., an edge) over and over again at different locations


Multiple Input Channels

- Input may have multiple input channels
 - Color image has 3 “channels” for red/green/blue
 - Input is actually $3 \times W \times H$
 - Solution: Kernel must be of size $C_{in} \times K \times K$
 - Where C_{in} is number of input channels




Multiple Input Channels

- Input may have multiple input channels
 - Color image has 3 “channels” for red/green/blue
 - Input is actually $3 \times W \times H$
 - Solution: Kernel must be of size $C_{in} \times K \times K$
 - Where C_{in} is number of input channels



Multiple Input Channels

- Input may have multiple input channels
 - Color image has 3 “channels” for red/green/blue
 - Input is actually $3 \times W \times H$
 - Solution: Kernel must be of size $C_{in} \times K \times K$
 - Where C_{in} is number of input channels



Multiple Output Channels

- What if you want more than one kernel?
 - Can have multiple kernels, each to detect a different thing
 - One for vertical lines, one for horizontal lines, etc.
 - So the total size of kernel tensor is $C_{\text{out}} \times C_{\text{in}} \times K \times K$

-1	2	-1
-1	2	-1
-1	2	-1


Kernel[0,0,:,:]

-1	-1	-1
2	2	2
-1	-1	-1

Kernel[1,0,:,:]

Stride and Padding

- Stride: As you slide across image, how big of a step do you take?
 - Default: stride=1 pixel
 - Can choose larger stride to reduce dimensionality
- Padding: Can pad the edges of images with 0's
 - For K=3 and no padding, width/height shrink by 2 each time
 - Adding width-1 padding on each side prevents this
 - For K=5, pad by 2, etc.
 - Default: No padding




Padding

Outline


- Regularizing Neural Networks
- Intuition for hierarchical features
- Extracting features with convolutions
- **Convolutional Neural Networks**

Convolutional Neural Networks (CNNs)




- How to incorporate convolutions into a full model?
- Basic idea: Use convolutions at beginning, then fully connected layer at end

Convolutional Layers




- First step: Convolutional Layer + ReLU
- Analogous to Linear layer + ReLU
 - Convolutional layer is just a special type of linear layer with local receptive fields & weight sharing!
 - So we again want to apply a non-linearity after the linear operation
- ReLU is standard for CNNs

Pooling




- Goal: Make receptive field bigger as we process the image
 - Early: Look for edges (small patch)
 - Later: Look for moose (larger patch)
- How do we do this? Pooling!
- Effectively we reduce resolution of input by a factor of P (often $P=2$)
 - Average pool: Average in each 2×2 patch
 - Max pool: Max in each 2×2 patch

More Conv + ReLU + Pool




- Can stack multiple Conv + ReLU + pool blocks
- Similar to increasing number of hidden layers in MLP
- Deeper layers convolutional layers have larger effective receptive field
 - Can learn higher-level concepts

Fully connected layers




- At the very end, we want fully global processing
- Fully connected layers are good at this!
- First flatten from [channels x width x height] to a flat vector
- Then do a MLP (e.g., 2-layer neural network) on top

Keeping the dimensions straight




- Suppose convolution kernels are 3×3 , 10 output channels, pooling is 2×2 , no padding, stride=1
 - Each convolution operation loses $3-1=2$ in width and height
- In code, also a “batch” dimension because we process all examples in batch together

How does backprop learn features?



- Every convolution & fully connected layer has (many) parameters
 - Convolutional: Kernel with $\text{\#outChannels} \times (\text{\#inChannels} \times K \times K + 1)$ params
 - Fully connected: $\text{\#outDimensions} \times (\text{\#inDimensions} + 1)$ params
- These all have to get learned by backprop + gradient descent on the loss


How does backprop learn features?



- Training example $(x^{(1)}, y^{(1)})$: $\partial(\text{Loss})/\partial(h) > 0$
 - Means that making h **smaller** leads to lower loss
- Training example $(x^{(2)}, y^{(2)})$: $\partial(\text{Loss})/\partial(h) < 0$
 - Means that making h **larger** leads to lower loss


- h is output of “classifier”
- Gradient tunes classifier parameters to make output larger on some examples, smaller on others

How does backprop learn features?



- Backpropagation: Does making c bigger change h in good or bad way?
- Sum up these considerations over all hidden units that depend on c
- Train convolutional kernel parameters so that value of c leads to [values of h 's that lead to good outputs]
- And so on for earlier layers...

What features do CNNs learn?



- Kernels of AlexNet first layer
 - Each one is 3 (for RGB) x 11 x 11
- What is learned?
 - Edge detectors in different directions and widths
 - Patches of various colors

What features do CNNs learn?



Each Row: Images that activate a different neuron in 5th POOL layer of AlexNet

Conclusion

- Convolution: Restricted linear operation parameterized by a small kernel
- Convolutional layers extract useful features for images
 - Motivation #1: Local Receptive Fields
 - Motivation #2: Weight Sharing
- Standard CNN architecture
 - Start: Convolutional layer + ReLU + Max Pooling
 - End: Fully connected layer

-1	2	-1
-1	2	-1
-1	2	-1

Kernel
(K=3)

0	0	0	0	0	0
0	1	0	0	0	0
0	1	1	1	1	1
0	1	0	0	0	0
0	0	0	0	0	0

Input

3	-1	0	0
5	-2	0	0
3	-1	0	0

Output